

# Lecture 7

COMP 3760

Data Structures and Graphs

Text chapter 1.4, 3.5, 5.3

# Fundamental Data Structures

(Chapter 1.4)

# Data Structures

- A *data structure* is a particular way of storing and organizing data
- Data structures and algorithms are often deeply interconnected
  - The way you organize data affects the performance of your algorithm
- We've *mostly* been using arrays ... so far

# Fundamental Data Structures

- Linear Data Structures
  - Array
  - Linked list
  - Stack
  - Queue
- Set
- Dictionary (Map)
- Tree
- Graph

# Arrays

- A sequence of  $n$  items of the same type, accessed by an index



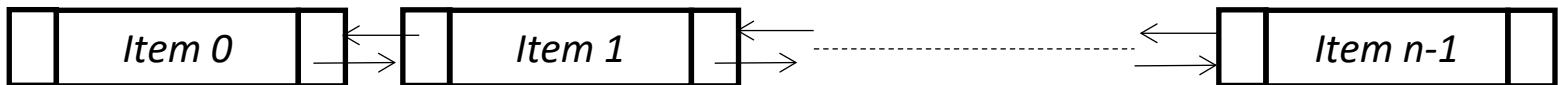
- The good:
  - Each item accessed in same constant time
- The bad:
  - Size is fixed
  - Insertion / deletion in an array is time consuming – all the elements following the inserted element must be shifted appropriately

# Linked Lists

- (singly) A sequence of zero or more elements called *nodes*, consisting of data and a pointer



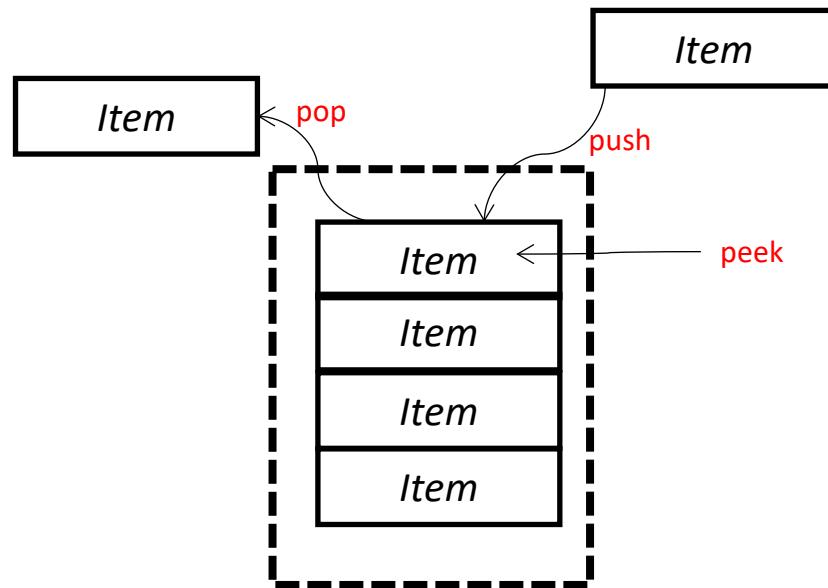
- (doubly) Pointers in each direction



# Linked Lists

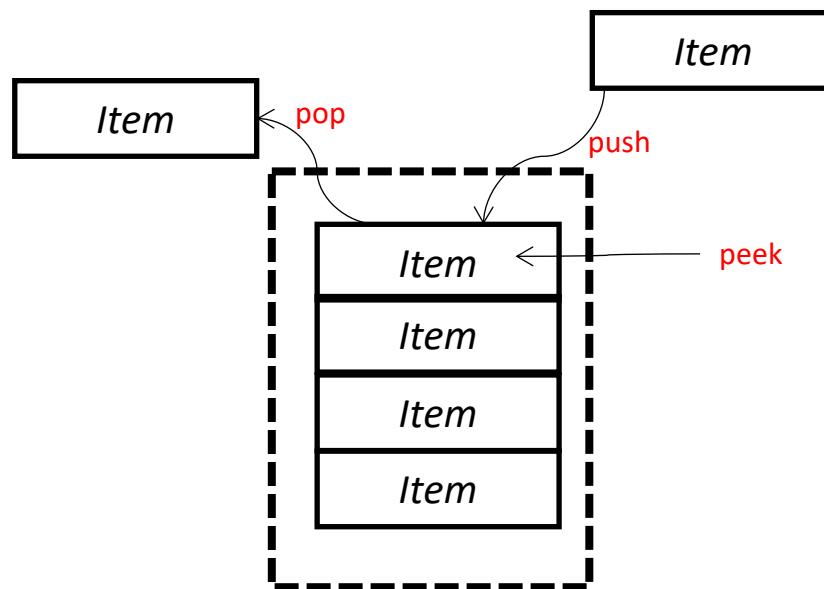
- Linked lists provide two key advantages over arrays
  - Dynamic size
  - Ease of insertion/deletion
- Linked lists have some drawbacks:
  - Random access is not allowed

# Do you know what this is?



# Stack

- Like a stack of plates
- Last-in-first-out (LIFO)



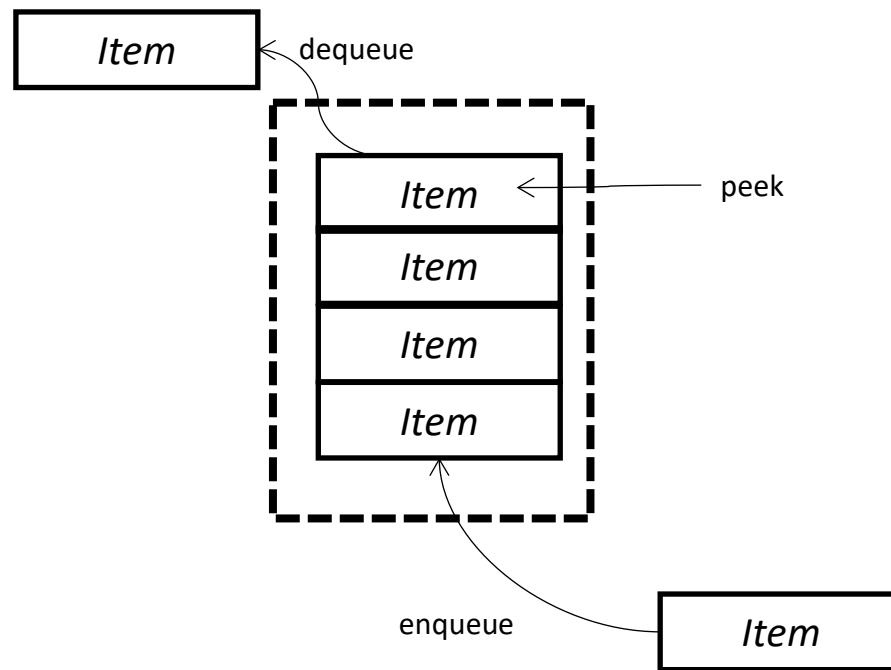
# Operations on a stack

- Insert operation is called Push
- Delete operation is called Pop
- Examining the top item is Peek
- Example application:
  - Analysis of languages (e.g. properly nested brackets)
  - Properly nested: (())
  - Wrongly nested: ())()

# Abstract Data Type

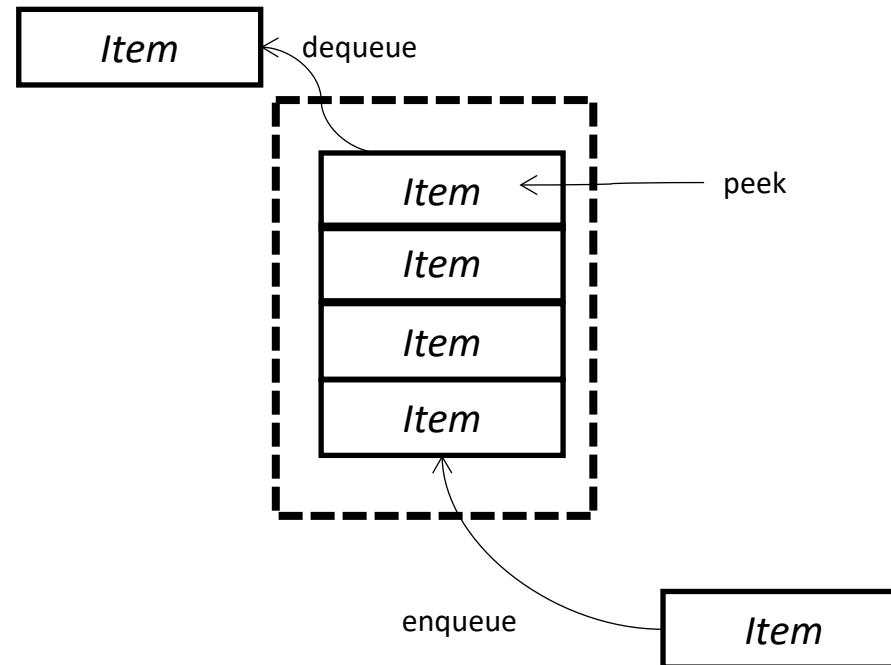
- Often a data structure is closely associated with a set of available operations
- **Data structure + operations = *abstract data type***
  - From an OOP perspective, think about members (methods) of a class
- Example 1: priority queue
  - Underlying implementation was a max-heap
  - Operations were Insert and deleteMax
- Example 2: stack
  - Operations are push, pop, peek

# How about this one?



# Queues

- Like a line-up
- First-in-first-out (FIFO)



# Operations on a queue

- Adding to the queue is Enqueue
- Removing from the queue is Dequeue
- The top/front element is the Head (sometimes there is a “Peek” method)

# Set

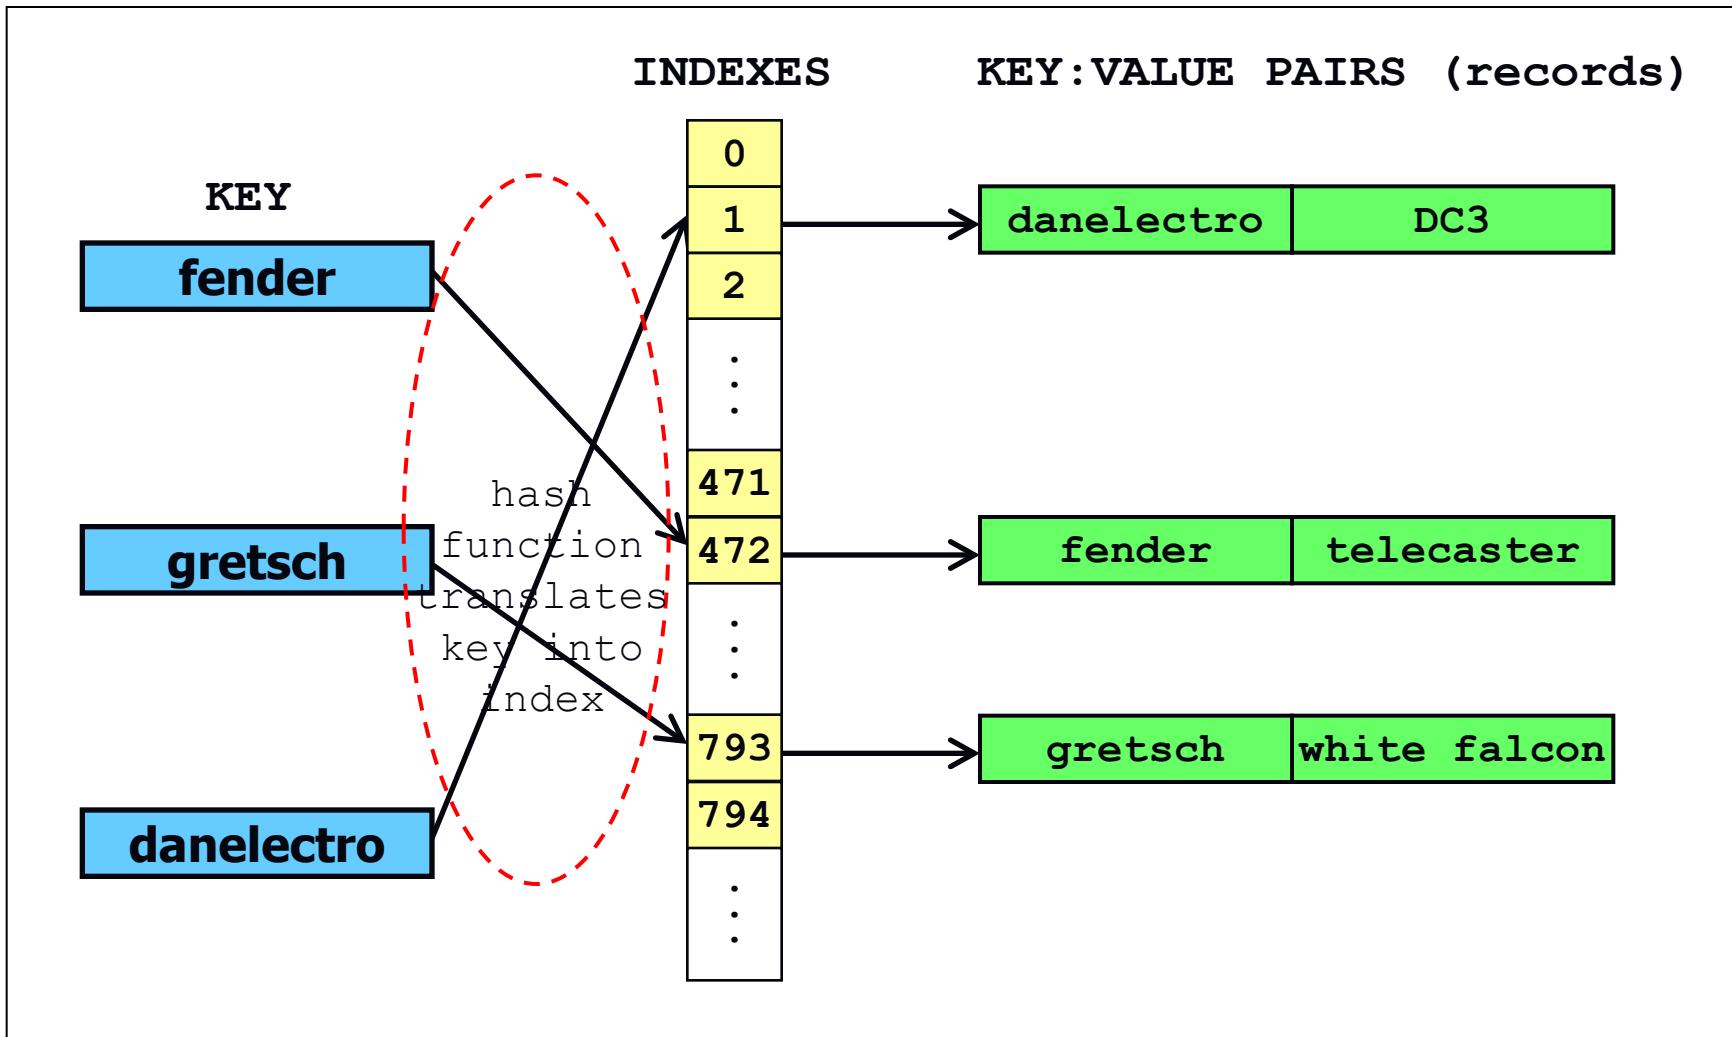
- Like a set in math, e.g. {1, 2, 3, 4}
  - Sets cannot contain duplicate items
- Operations on a Set:
  - Add an item to it
  - Remove an item from it
  - Check if an item is in it
  - Iterate over it (loop based on all items, one-by-one)

# Set in Java

- Different ways to implement a set
  - HashSet
    - Faster implementation, but it is unordered
  - TreeSet
    - Slower, but the items are available in order

# Map (as a hash table)

- A **Map** is a lookup table that takes a **key** and returns a **value**
  - the most common implementation is as a hashtable (hashmap)

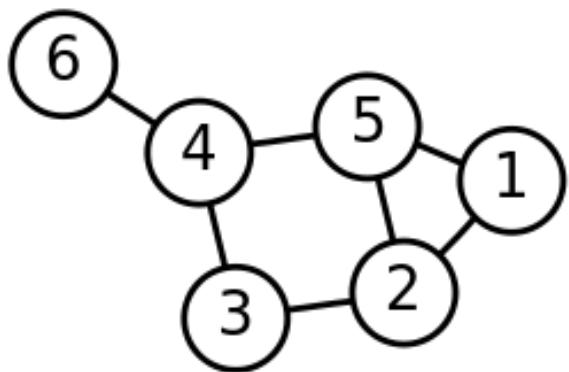


# Graphs

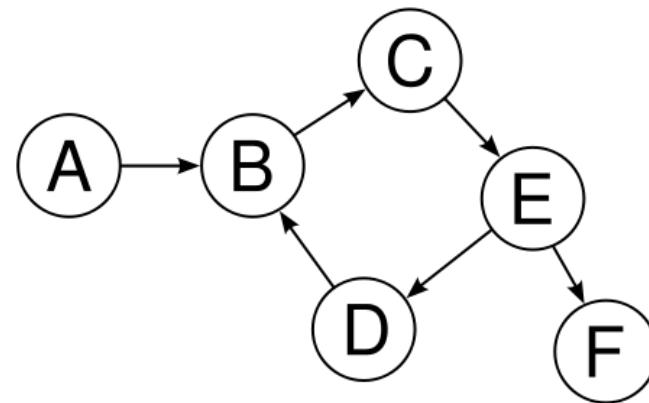
(Still in Chapter 1.4)

# Graphs

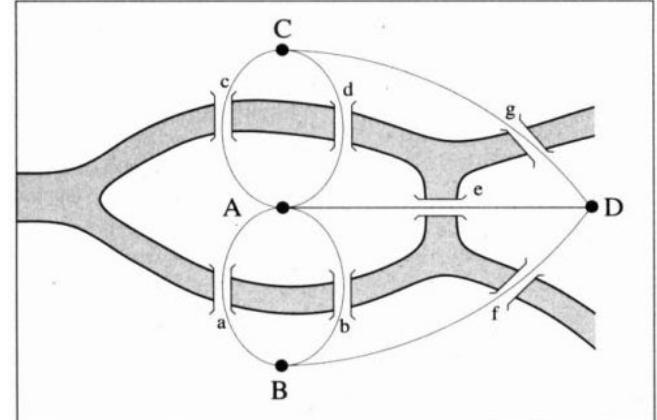
- $G = (V, E)$ 
  - $V$  is a set of *vertices*
  - $E$  is a set of *edges*



Undirected



Directed



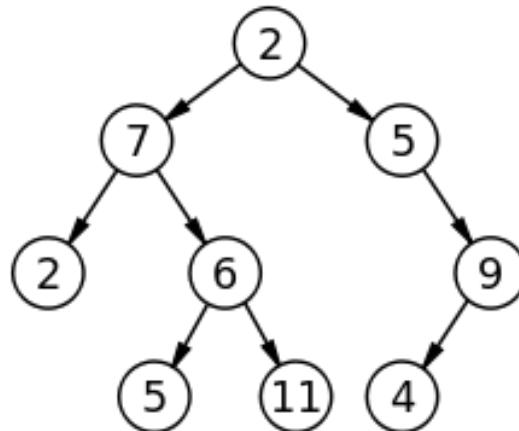
Motivation: Real world connections

# Some special graphs

- *Connected graph*
  - A graph where there is a path available between any two vertices
- *Cyclic graph*
  - A graph containing at least one cycle
- *Acyclic graph*
  - A graph containing no cycles
- *Tree*
  - Any connected + acyclic graph
- *Complete graph*
  - Every pair of vertices is connected by an edge
- *Weighted graph*
  - Every edge has an associated value

# Trees

- Connected, acyclic graphs
  - Usually we think of trees as having a root
- Representing data in a tree can speed up your algorithms in many natural problems

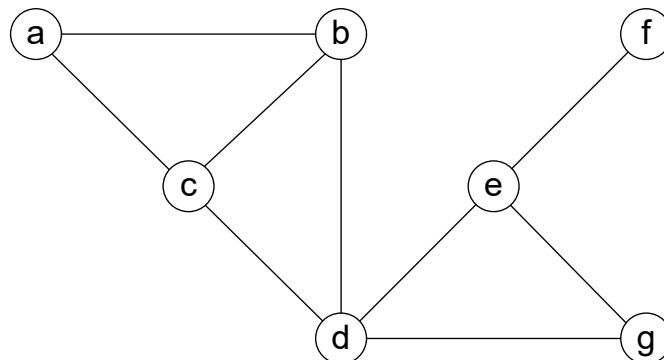


# Representation of graphs

- Two common ways to represent graphs:
  - Adjacency matrix
    - $|V| \times |V|$  matrix
    - Cell  $i, j$  represents an edge from vertex  $i$  to  $j$
  - Adjacency lists
    - $|V|$  linked lists – one for each vertex, showing all the neighbours of that vertex

# Representation: Adjacency Matrix

- ▶ For this graph:

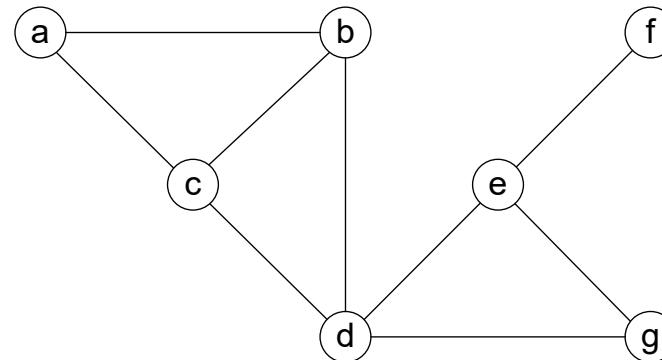


- Adjacency matrix is the following:

	a	b	c	d	e	f	g
a	0	1	1	0	0	0	0
b	1	0	1	1	0	0	0
c	1	1	0	1	0	0	0
d	0	1	1	0	1	0	1
e	0	0	0	1	0	1	1
f	0	0	0	0	1	0	0
g	0	0	0	1	1	0	0

# Representation: Adjacency List

- ▶ For the same graph:



- Adjacency list is the following:

- For vertex a: [b] → [c]
- For vertex b: [a] → [c] → [d]
- For vertex c: [a] → [b] → [d]
- For vertex d: [b] → [c] → [e] → [g]
- For vertex e: [d] → [f] → [g]
- For vertex f: [e]
- For vertex g: [d] → [e]

# Representing Graphs

## 1. Adjacency matrix

- Or Weight Matrix for weighted graphs

## 2. Adjacency lists

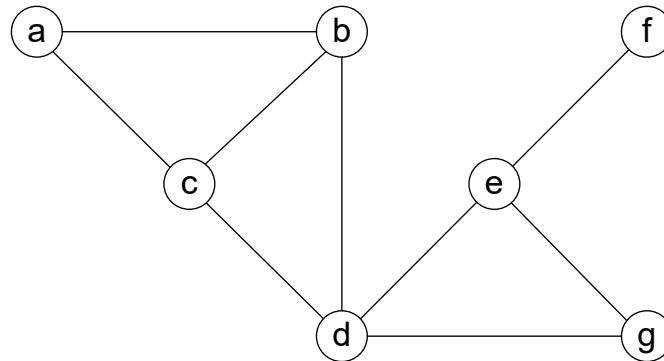
- A list of vertices connected to each vertex
- Which one to use?
  - Depends on the nature of the graph (sparse or not)
  - Depends on the algorithm

# Graph Algorithms

(Chapter 3.5)

# Graph Traversal

- Many real-world problems require processing of each vertex (or edge) in a graph



- Routing a message on a network
- Web crawling
- Social networking
- Garbage collection
- Solving puzzles

# Graph Traversal Algorithms

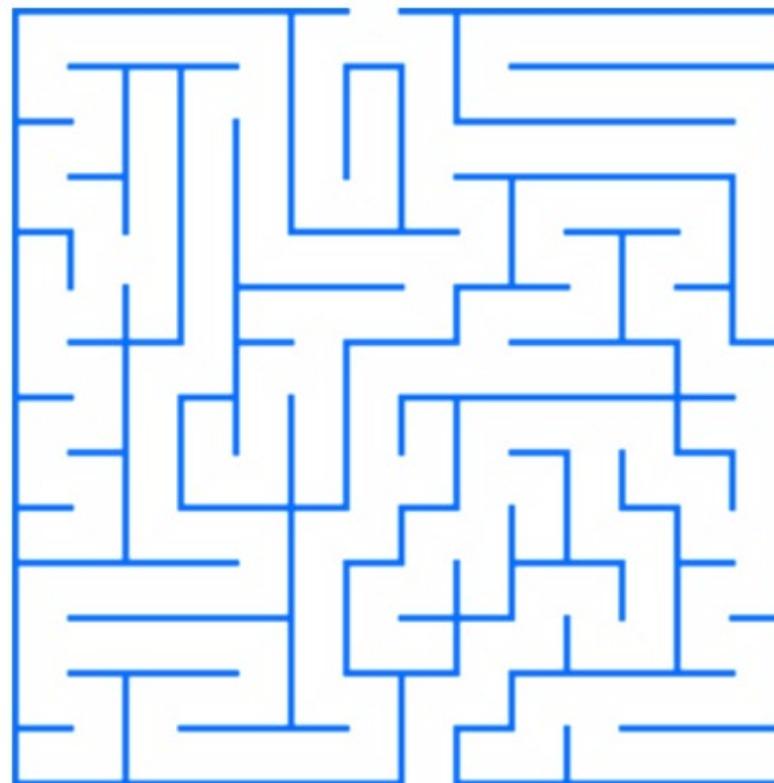
- Graph traversal algorithms give a method for *systematically processing* all vertices

Basic idea: "visit" all the vertices, one at a time,  
*marking* them as we visit them

- Two approaches:
  - Depth-First Search (DFS)
  - Breadth-First Search (BFS)

# Depth-first search (DFS)

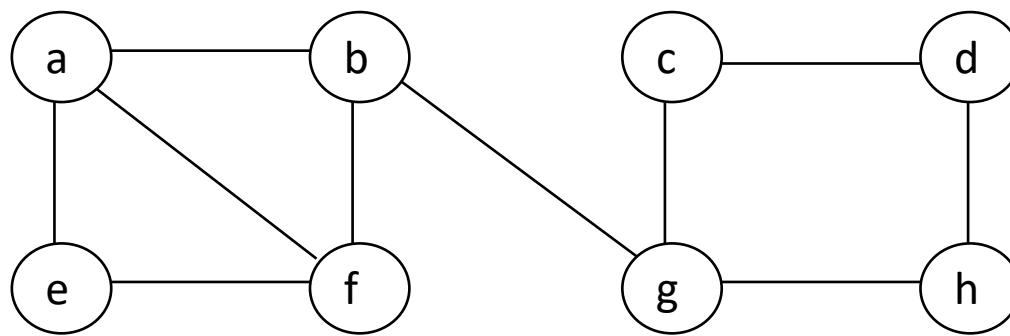
- Think about how you might try to find your way through a maze



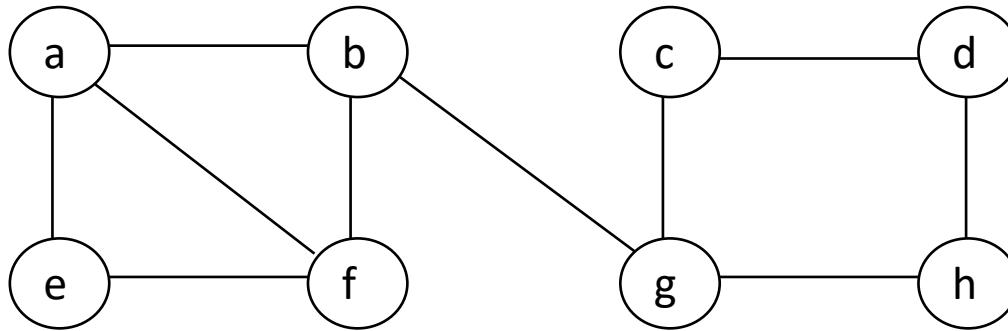
# Depth-first search (DFS)

- Visits all vertices by always *moving away* from the last vertex visited (if possible)
  - Backtracks at “dead-ends” (no adjacent, unvisited vertices)
- Implementation often uses a stack of vertices being processed
- Follows a tree-like route throughout the graph

# DFS example



# DFS example



Backtrack/finish order: e f h d c g b a

DFS order: a b f e g c d h

# Some notes on DFS

- To track the progress of the algorithm we use a stack
  - When we make a recursive call, e.g.  $\text{dfs}(v)$ , we push  $v$  onto the stack
  - When  $v$  is a dead-end (i.e. no more neighbors to visit) it is popped off the stack
- Our convention: break ties for “next neighbor” by using some natural order
- Typical results from running DFS can be:
  - List of vertices in order visited
  - List of vertices in order of “dead-ends” (when popped from stack)
  - DFS Tree – tree containing all the edges that were used to visit nodes
    - Unused edges of  $G$  (edges not in DFS tree) are called “back edges”

# DFS algorithm

```
Algorithm Depth_First_Search(Graph G)
    // Graph G = {V,E}
    initialize visited to false for all vertices
    for each vertex v in V
        if v has not been visited
            dfs_helper(v)

function dfs_helper(Vertex v)
    visit node v
    for each vertex w in V adjacent to v
        if w has not been visited
            dfs_helper(w)
```

- “Visit node v” means doing whatever you need to do at each node

# Common uses of DFS

- Find a spanning tree of a graph
- Find a path between two vertices  $v$  and  $u$
- Find a path out of a maze
- Determine if a graph has a cycle
- Find all the connected components of a graph
- Search the state-space of problems for a solution  
(AI)
- Many more!

# Efficiency of DFS

- The basic operation is the if statement in `dfs_helper()`:

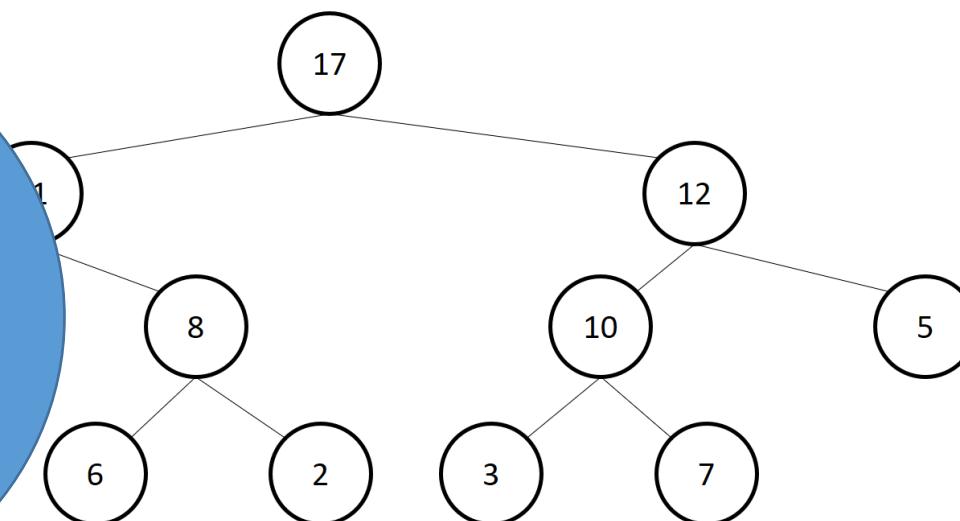
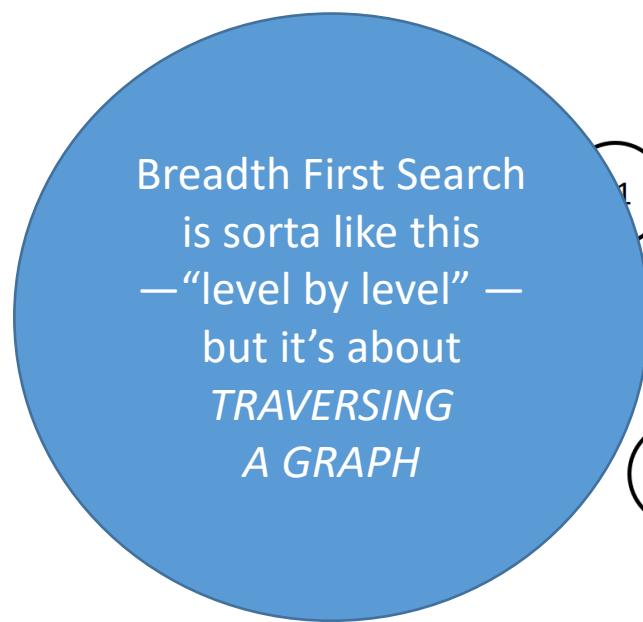
```
for each vertex w in V adjacent to v
    if w has not been visited
        dfs(w)
```
- Each time `dfs_helper()` is called, this loop examines all the neighbors of some vertex  $v$  ... eventually it looks at ALL the neighbors of ALL the vertices
  - Therefore the number of basic operations depends on the data structure used to implement the graph
- Basically we need to visit each element of the data structure exactly once. So the efficiency must be:
  - $O(|V|^2)$  for adjacency matrix
  - $O(|V|+|E|)$  for adjacency lists

# Which is better/worse?

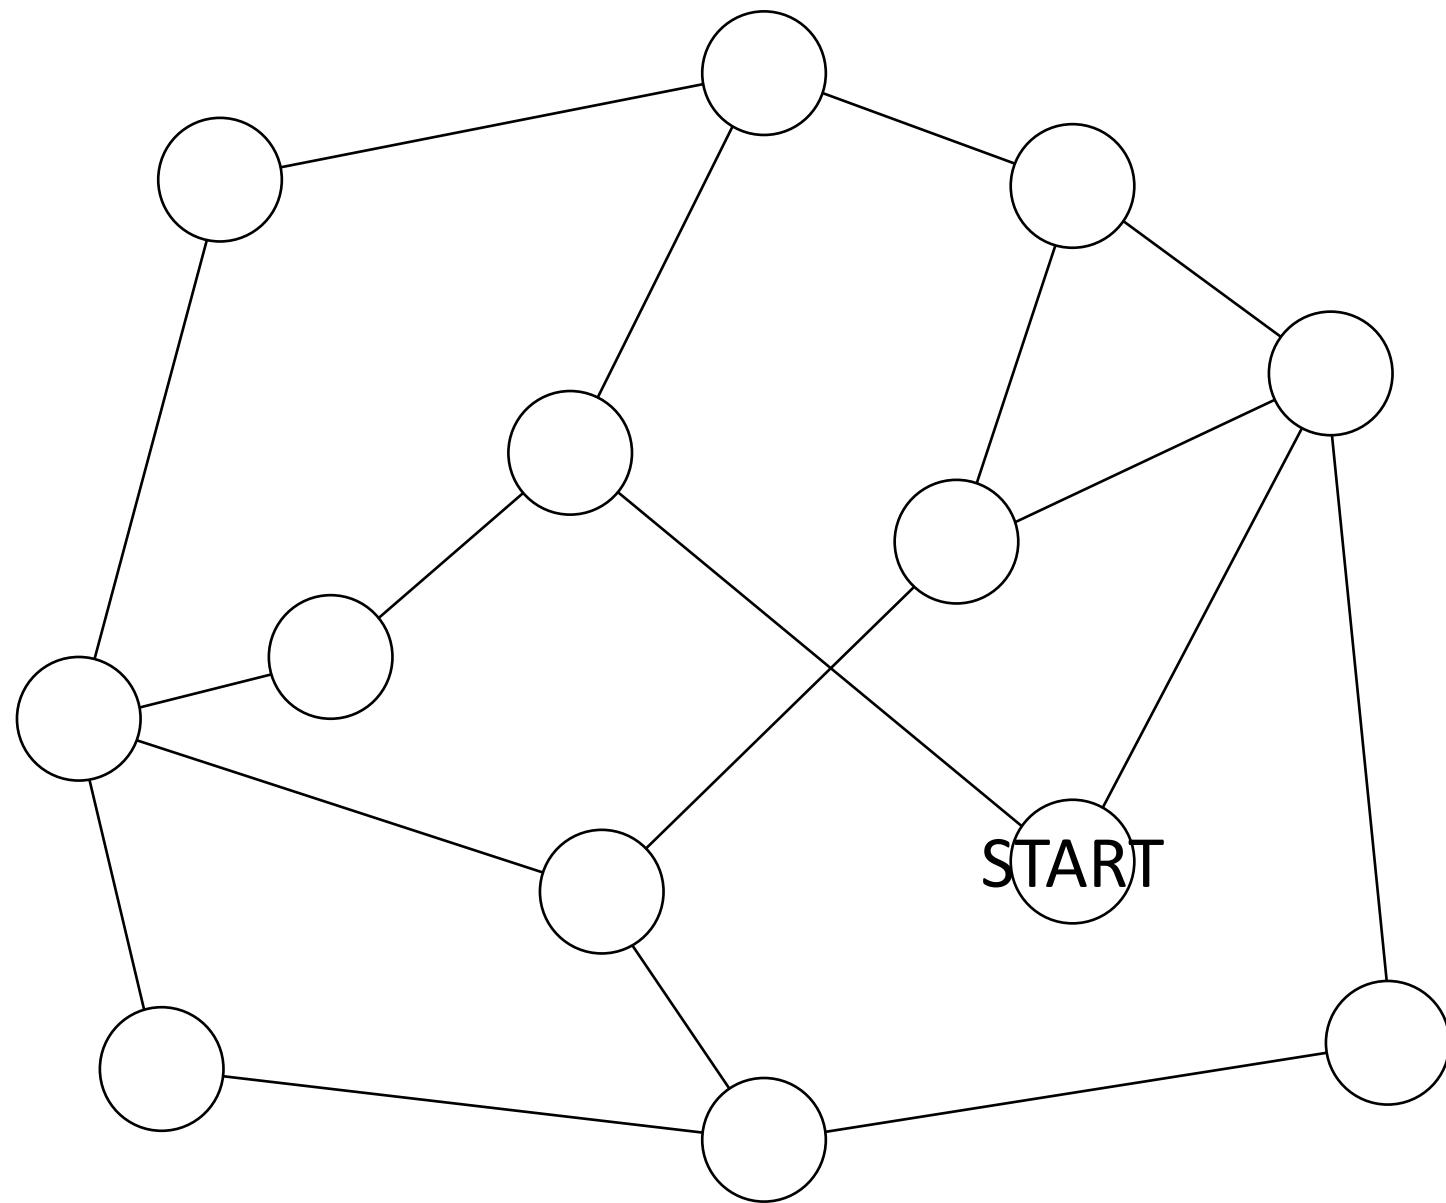
- $O(|V|^2)$
- $O(|V|+|E|)$

# Breadth-first search (BFS)

- Recall the “trick” where we used an array to store a (very specific type of) binary tree (a heap).



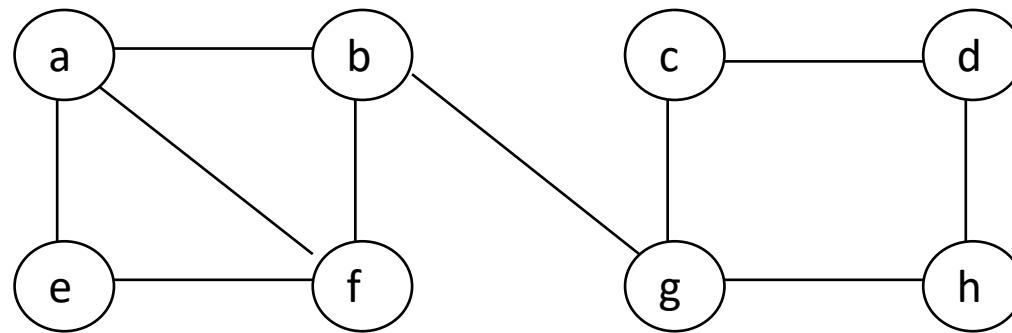
17 11 12 9 8 10 5 1 4 6 2 3 7



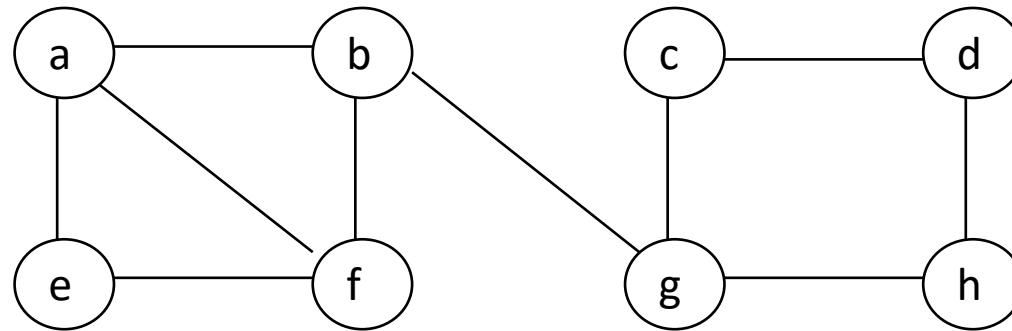
# Breadth-first search (BFS)

- Visit all neighbors “the same distance” from starting vertex
  - Visit immediate neighbors first
  - Then the neighbors of all those vertices
  - Etc.
- Instead of a stack, BFS uses a queue
- Follows a tree-like route throughout the graph, but perhaps a different tree than DFS

# BFS example



# BFS example



BFS order: a b e f g c h d

# BFS algorithm

```
Algorithm Breadth_First_Search(Graph G)
// Graph G = {V,E}
    initialize visited to false for all vertices
    for each vertex v in V
        if v has not been visited
            bfs_helper(v)

function bfs_helper(Vertex v)
    visit node v
    initialize a queue Q
    add v to Q
    while Q is not empty
        for each w adjacent to Q.head
            if w has not been visited
                visit node w
                add w to Q
    Q.dequeue()
```

- Uses a queue (FIFO) to determine which vertex to visit next
- Edges that are in G but not in the resulting BFS tree are called “cross-edges”

# Notes on BFS

- Same efficiency as DFS:
  - Adjacency matrix:  $O(|V|^2)$
  - Adjacency list:  $O(|V|+|E|)$
- Yields just one ordering of vertices (order added/deleted from queue is the same)
  - Whereas with DFS, the order that vertices are *visited* may be different from the order they get *finished* (become dead-ends)

# BFS applications

- Really the same as DFS
- Sometimes one or the other may be better for specific problems

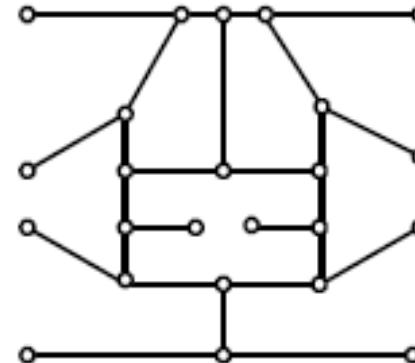
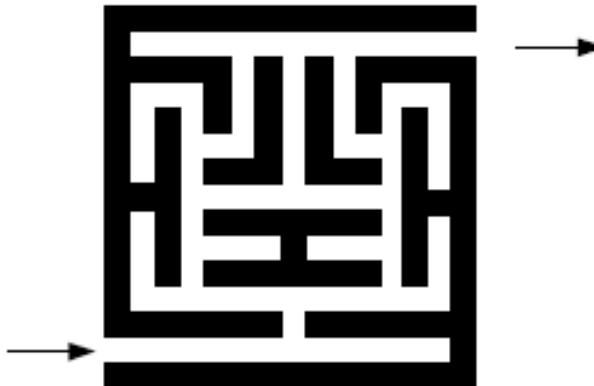
Some problems solvable  
with graph traversal

# Problem 1: Spanning Tree

- Given a connected graph  $G$ , find a spanning tree  $T$ 
  - This is a straight-up application of BFS (or DFS)
  - Build a new graph (the spanning tree) as we go
    - Initialize a new graph  $T$
    - Each time we visit a vertex, add the edge we used to  $T$
- BFS or DFS?
  - BFS usually gives shorter paths between vertices

# Problem 2: Solving a Maze

- Represent maze as a graph
  - Nodes for start, finish, intersections, and dead-ends
  - Find a path from *start* to *finish*
- BFS or DFS?
  - If interested in end-result:
    - BFS will find the shortest total path
  - If actually walking while solving:
    - DFS tends to result in less actual walking
    - BFS backtracks to parent nodes too often



# Problem 3: Shortest Path

- Find the shortest path between two vertices  $u$  and  $w$
- BFS or DFS?

# Problem 3: Shortest Path

- Find the shortest path between two vertices  $u$  and  $w$
- BFS or DFS?
  - BFS will find a shortest path
  - DFS will find a path – maybe not the shortest one
- Idea of algorithm (and why it works):
  - First, use  $\text{bfs}(u)$  to create a spanning tree  $T$  with  $u$  as the root. Note that all paths that appear in  $T$  are the shortest paths from  $u$  to their respective vertices
  - Then, use DFS on  $T$  to find a path from  $u$  to  $w$  (as in the maze problem)

# Problem 4: Determine Connectivity

- Determine if a graph is connected
- BFS or DFS?
  - Either will work
  - Think about this yourself!
  - What modification(s) do you need to make to the algorithm to answer this question?

# Graph Algorithms: Binary tree traversal

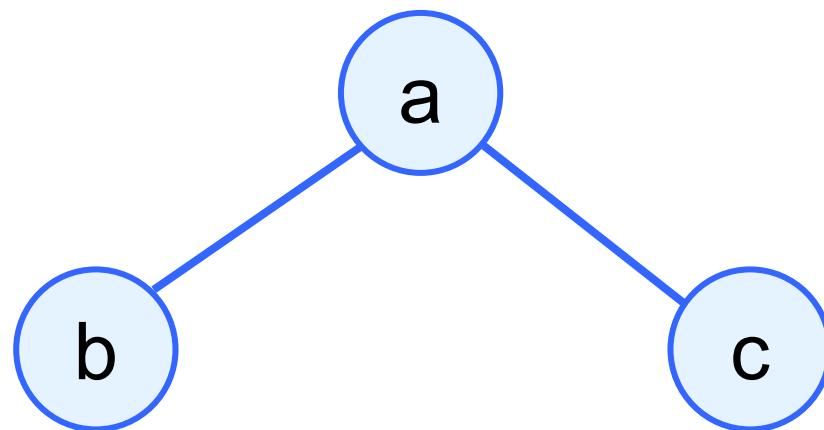
Textbook: Chapter 5.3

# Tree traversal

- Traversing a tree means to visit all of the nodes of the tree
- We've already seen DFS and BFS (for graphs)
- Here are a few traversals specific to *binary trees*:
  - Preorder – root *before* the children
  - Inorder – root *between* the children
  - Postorder – root *after* the children

# Preorder traversal

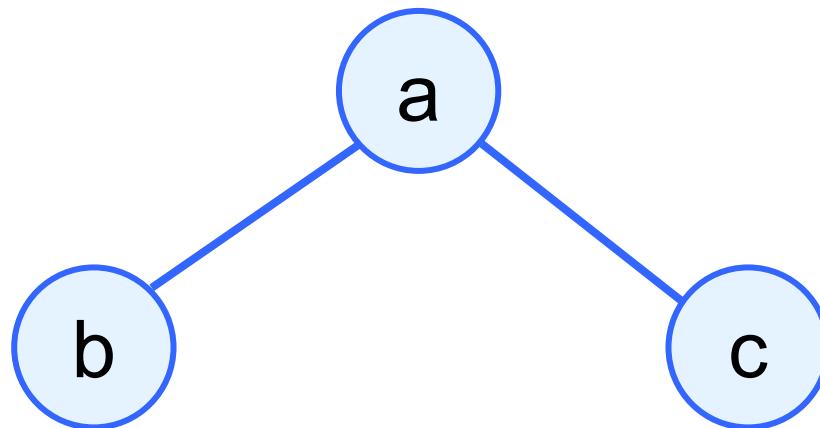
1. Visit the root
2. Traverse the left subtree
3. Traverse the right subtree



Preorder traversal is: a b c

# Inorder traversal

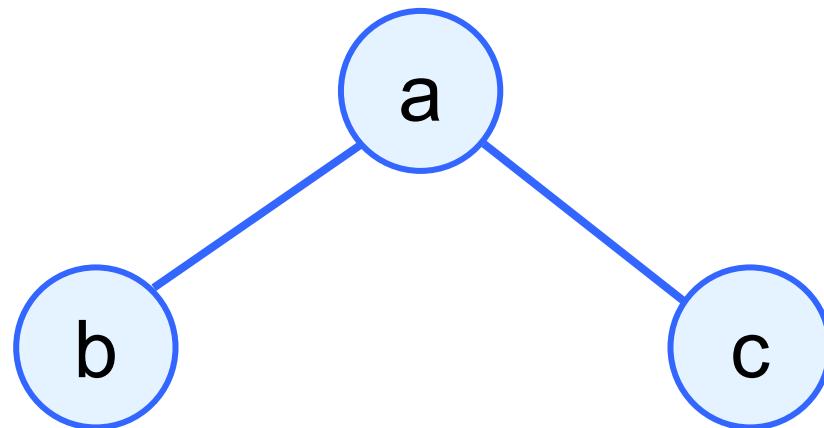
1. Traverse the left subtree
2. Visit the root
3. Traverse the right subtree



Inorder traversal is: b a c

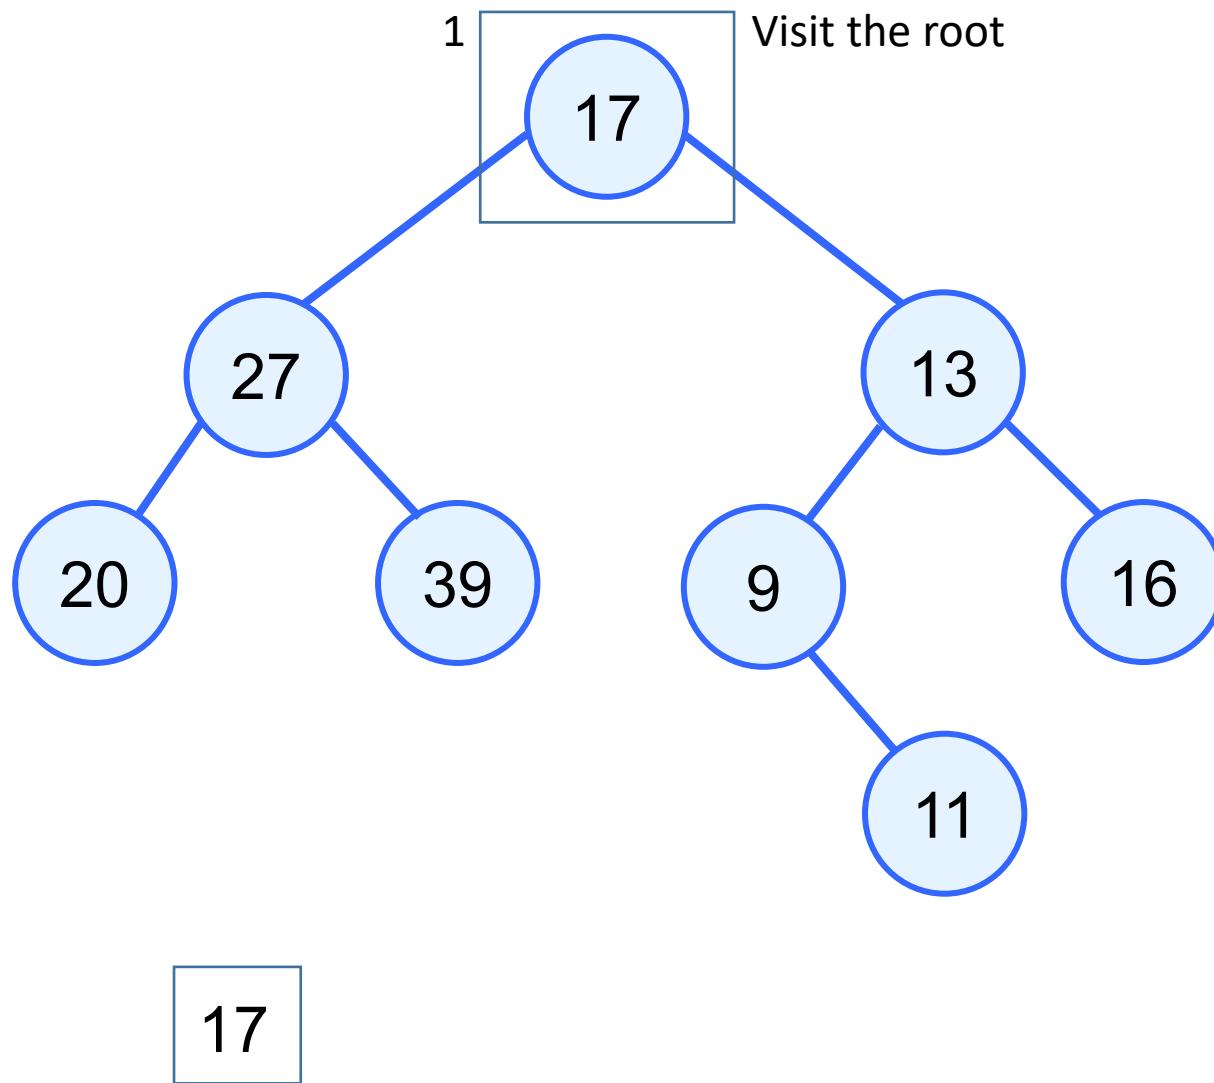
# Postorder traversal

1. Traverse the left subtree
2. Traverse the right subtree
3. Visit the root

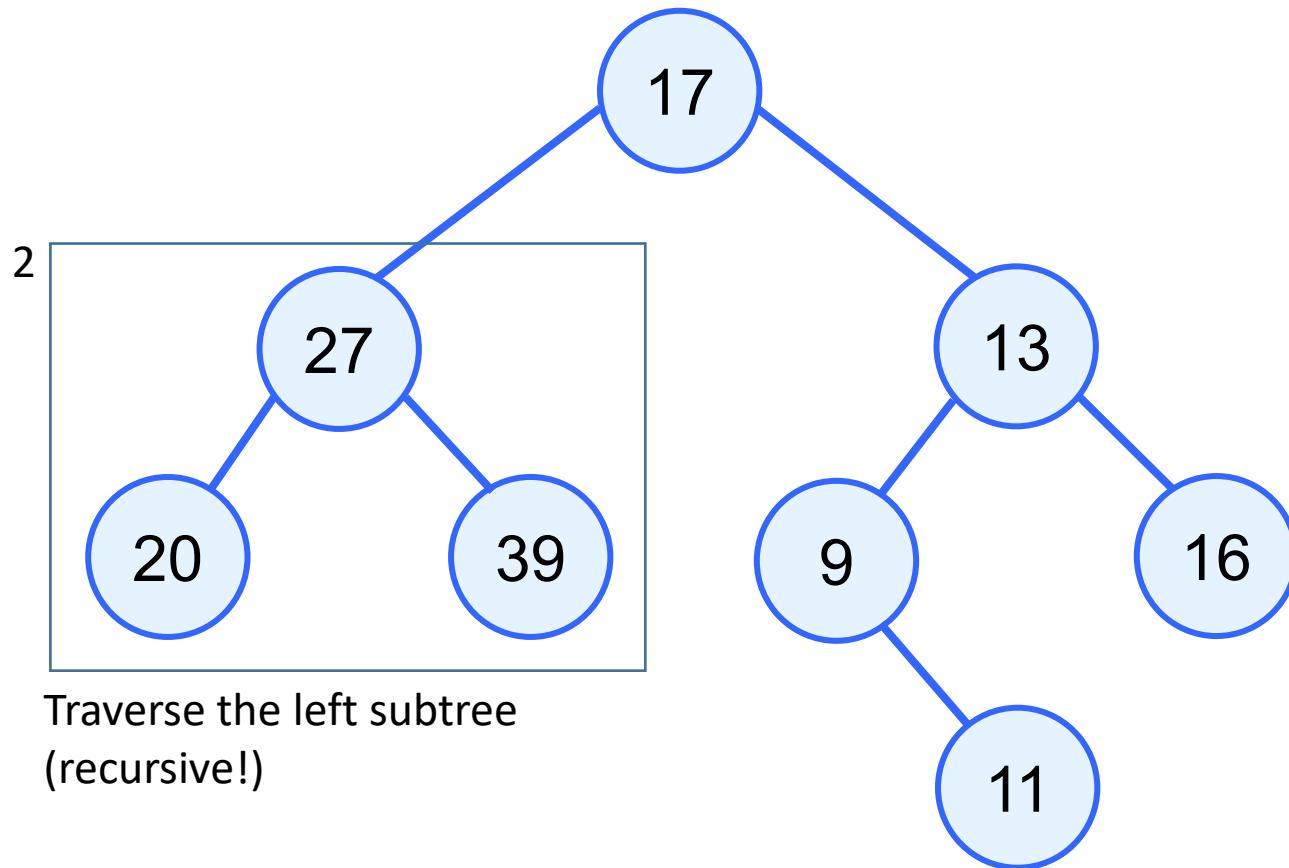


Postorder traversal is: b c a

# Preorder example



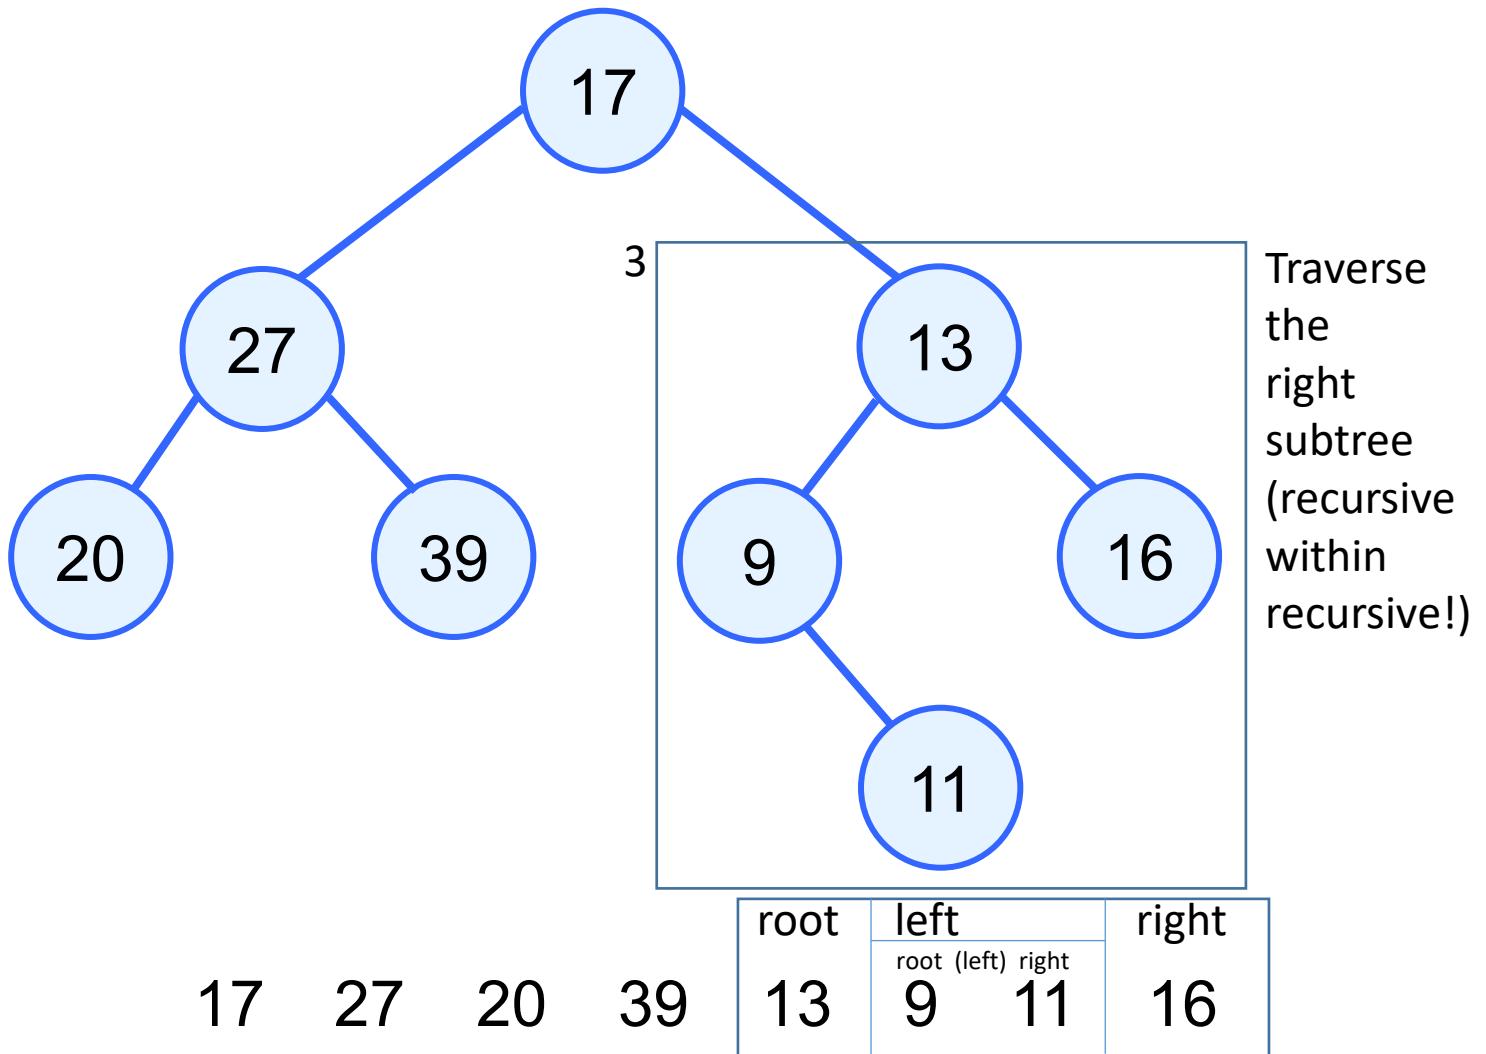
# Preorder example



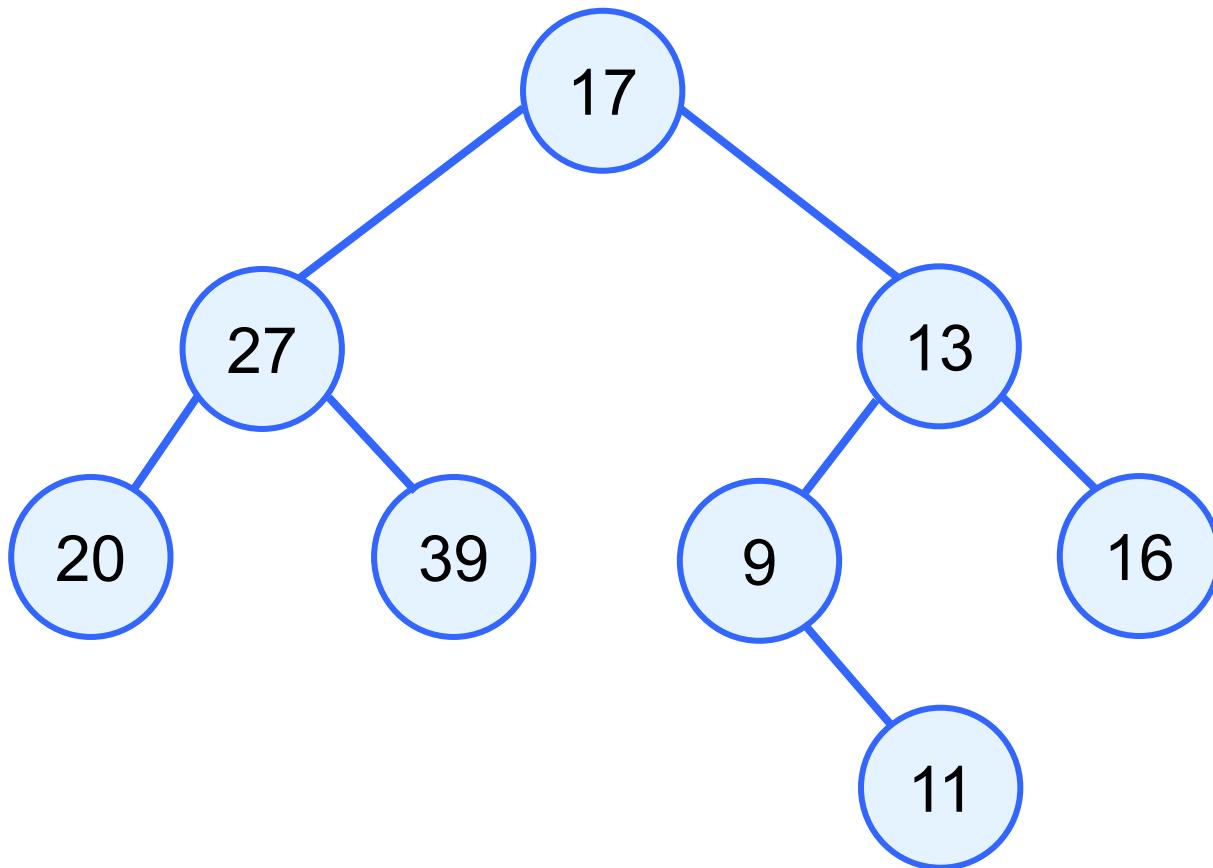
Traverse the left subtree  
(recursive!)

root	left	right
17	27	39

# Preorder example



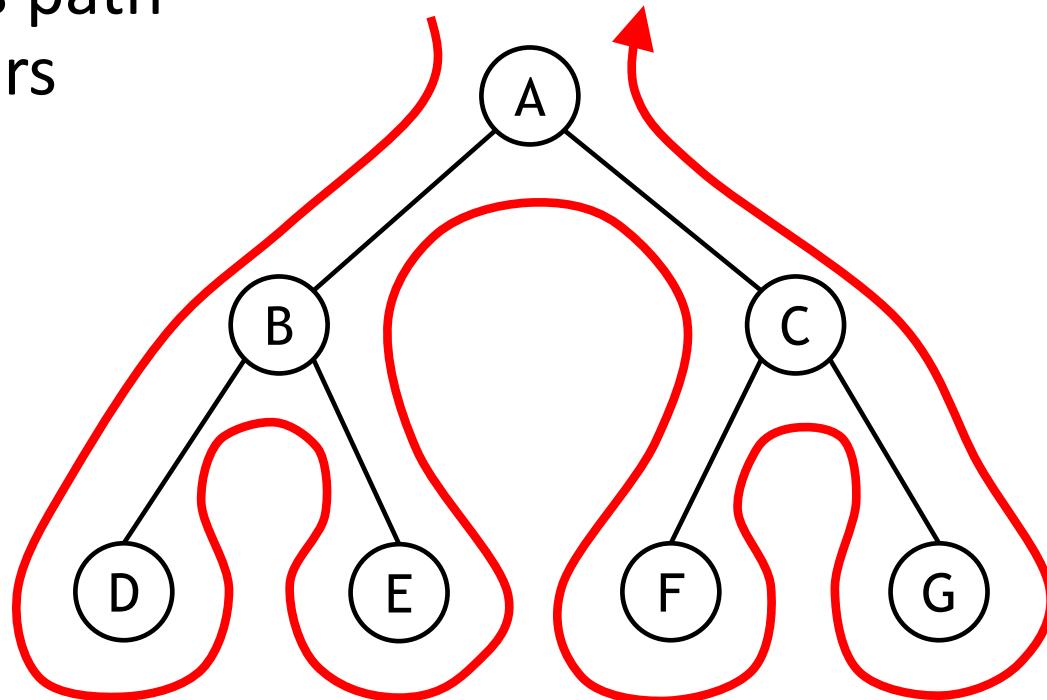
# Preorder example



17 27 20 39 13 9 11 16

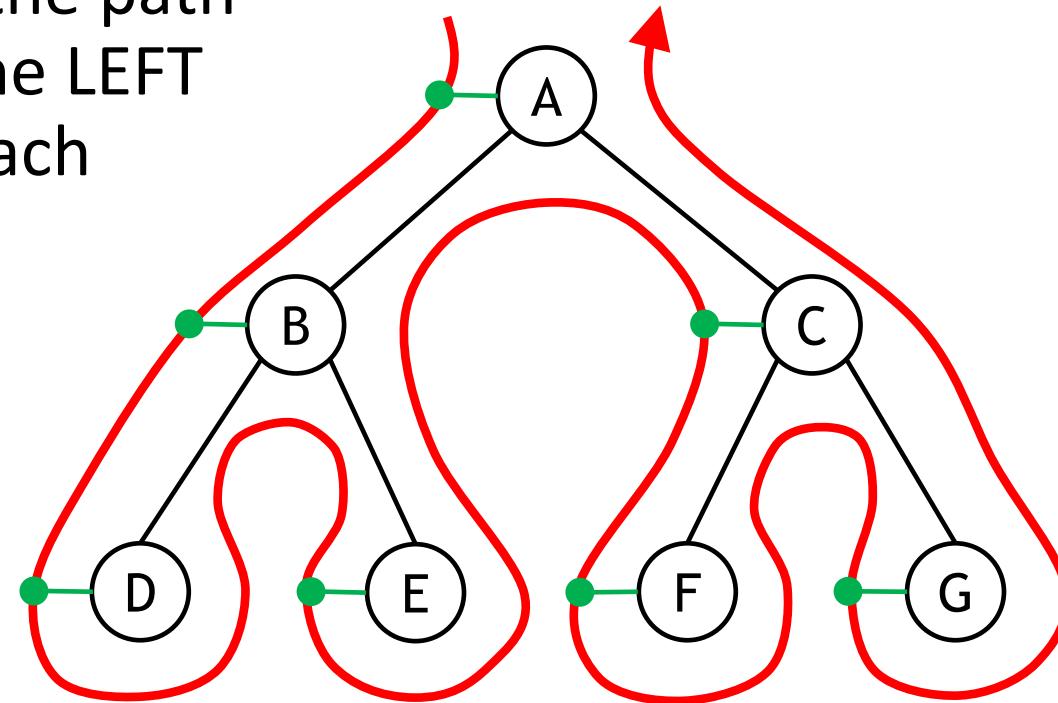
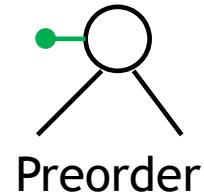
# Another way to think about it

- Consider this path that meanders past all of the nodes
- ALL three traversals follow this path!



# Preorder traversal

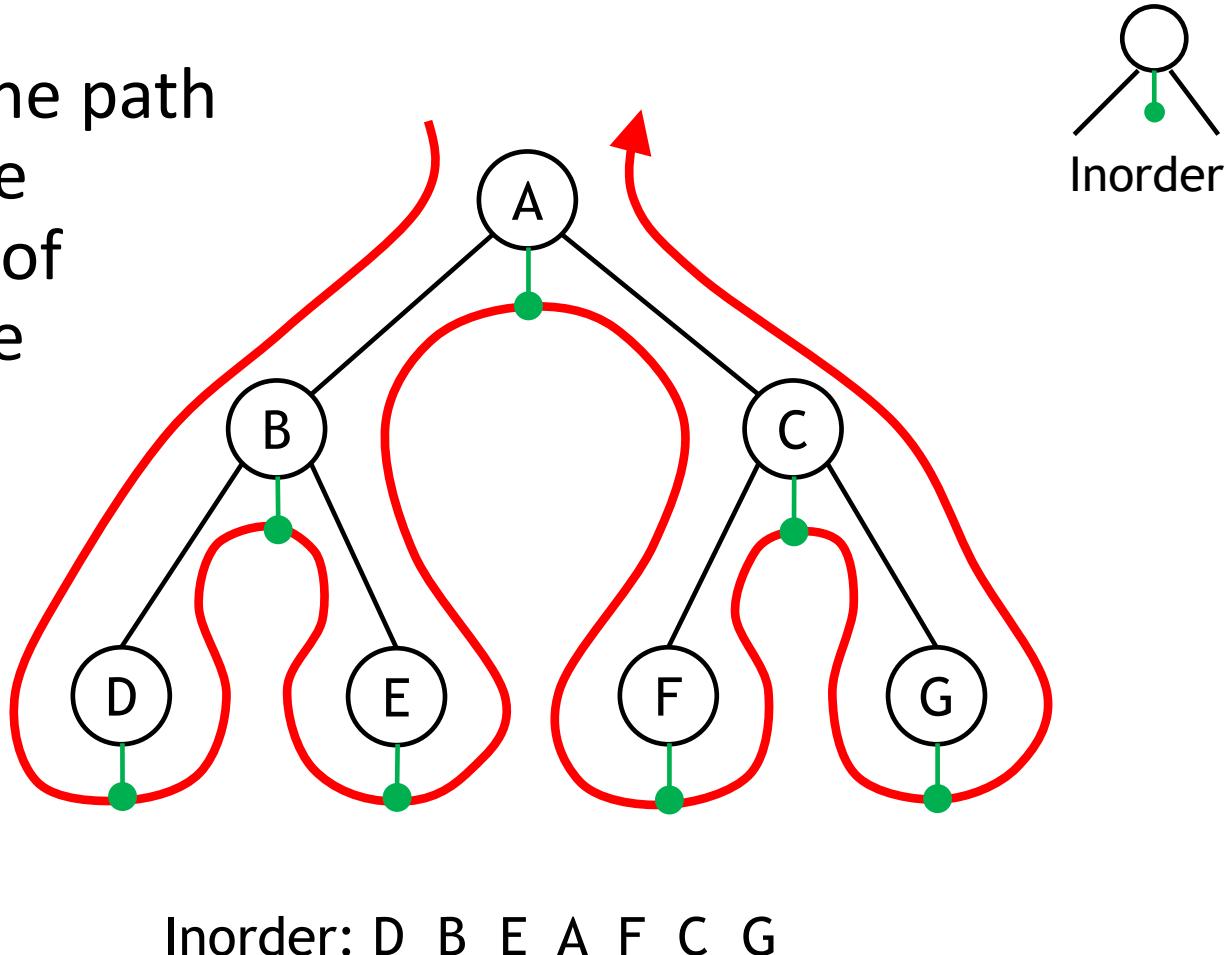
- Is when the path passes the LEFT side of each node



Preorder: A B D E C F G

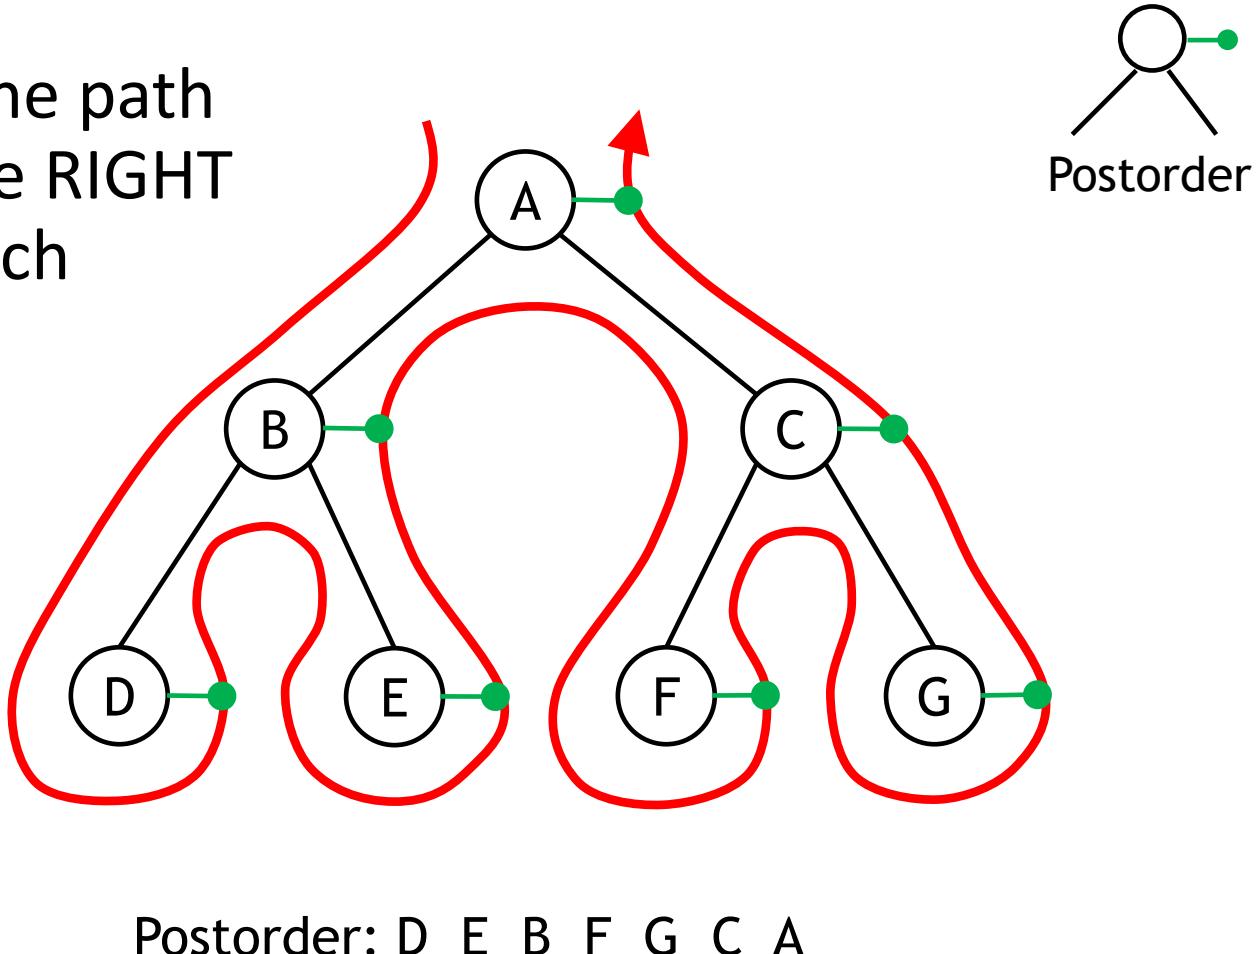
# Inorder traversal

- Is when the path passes the BOTTOM of each node



# Postorder traversal

- Is when the path passes the **RIGHT** side of each node



# Pseudocode

```
Algorithm preOrder (Node N)
if N != null
    Print N.value
    preOrder(N.leftChild)
    preOrder(N.rightChild)
```

```
Algorithm postOrder (Node N)
if N != null
    postOrder(N.leftChild)
    postOrder(N.rightChild)
    Print N.value
```

```
Algorithm inOrder (Node N)
if N != null
    inOrder(N.leftChild)
    Print N.value
    inOrder(N.rightChild)
```

# What if I told you

Preorder → A B D E C F G

Inorder → D B E A F C G

# Fun facts about pre/in/postorder

- Given pre + in, you can reconstruct the tree
  - (and also determine postorder)
- Given post + in, you can reconstruct the tree
  - (and also determine preorder)
- Given pre + post, you can only *sometimes* reconstruct the tree
  - For you to ponder: under what condition(s)?

# DFS algorithm

```
Algorithm Depth_First_Search(Graph G)
    // Graph G = {V,E}
    initialize visited to false for all vertices
    for each vertex v in V
        if v has not been visited
            dfs_helper(v)

function dfs_helper(Vertex v)
    visit node v
    for each vertex w in V adjacent to v
        if w has not been visited
            dfs_helper(w)
```

- “Visit node v” means doing whatever you need to do at each node

# BFS algorithm

```
Algorithm Breadth_First_Search(Graph G)
// Graph G = {V,E}
    initialize visited to false for all vertices
    for each vertex v in V
        if v has not been visited
            bfs_helper(v)

function bfs_helper(Vertex v)
    visit node v
    initialize a queue Q
    add v to Q
    while Q is not empty
        for each w adjacent to Q.head
            if w has not been visited
                visit node w
                add w to Q
    Q.dequeue()
```

- Uses a queue (FIFO) to determine which vertex to visit next
- Edges that are in G but not in the resulting BFS tree are called “cross-edges”