

4 - Discrete Probability Distributions

Example Suppose you roll five six-sided dice. Let X = the *sum* of the five dice. Events specified in terms of X include:


$$X = 30$$

$$X \geq 28$$



Example Let X = the number of CST graduates in a random sample of 6 who know how to construct a linked list in C++. Suppose the probability distribution of X is given by the following table and/or probability histogram.

x	$P(x)$
0	0.2620
1	0.3930
2	0.2459
3	0.0816
4	0.0160
5	0.0015
6	0+



What is $P(X \geq 3)$?

Definition A *discrete random variable* X is a random variable that has either a *finite* number of values or a *countable* number of values.

e.g., X = the number of children a random person has in their lifetime

Definition The *discrete probability distribution* of a random variable X tells us $P(X = x)$ for any possible value x . It can be given by:

- a table/histogram
- a formula

Requirements for a Discrete Probability Distribution

For any discrete random variable X , the following must be true about the probabilities $P(x)$:

$$\sum P(x) = 1 \quad \text{sum of all probabilities equals 1}$$

$$0 \leq P(x) \leq 1 \quad \text{for each possible value } x$$

4.1 - Mean and Variance of a Random Variable

Example What is the mean value of X = number rolled on a fair 6-sided die? Imagine many, many rolls:

```
> X.vals <- sample( 1:6, 100, replace=TRUE)
> X.vals
[1] 5 3 3 4 4 5 4 4 5 3 3 6 2 1 5 5 1 4 2 2 5 5 3 3 1 6 4 5 3 3 4
[32] 3 3 6 4 1 5 2 1 2 1 5 1 2 5 1 1 5 1 1 3 4 2 2 3 5 1 2 3 3 5 3
[63] 6 1 6 1 6 2 2 5 5 5 6 2 5 3 5 1 5 5 5 1 5 2 3 3 1 6 2 5 6 4 1
[94] 2 1 3 4 4 4 6
```

For these 100 simulated values, the mean \bar{X} is:

```
> sum(X.vals) / 100
[1] 3.38
```

If we could compute the mean for *all possible* rolls of the 6-sided die, we would get:

Definition If X is a discrete random variable with probability distribution given by $P(x)$ then we define the mean and standard deviation as follows:

$$\mu = \sum_x [x \cdot P(x)]$$

The mean μ is also called the “expected value” of X and can be written as $E[X]$.

$$\sigma^2 = \sum_x [(x - \mu)^2 \cdot P(x)]$$

$$\sigma^2 = \sum_x [x^2 \cdot P(x)] - \mu^2$$

Example For the CST graduates example above, we have:

x	$P(x)$	$x \cdot P(x)$	x^2	$x^2 \cdot P(x)$
0	0.2620	0.0000	0	0.0000
1	0.3390	0.3930	1	0.3930
2	0.2459	0.4918	4	0.9836
3	0.0816	0.2448	9	0.7344
4	0.0160	0.0640	16	0.2560
5	0.0015	0.0075	25	0.0375
6	0.0000	0.0000	36	0.0000
Total	1.000	1.2011		2.4045

$$\mu = 1.20 \text{ CST graduates}$$

$$\sigma^2 = 2.4045 - 1.2011^2 = 0.9619$$

$$\sigma = 0.98 \text{ CST graduates}$$

Example Suppose I ask: “Pick a random number between 1 and 100.” What probability distribution am I likely thinking of? What is the mean, and what is the standard deviation?

Example (Minimum of Two Dice) The random experiment is rolling two fair six-sided die. Let X = the *minimum* of the two die values. What is the expected value of X ? What is the standard deviation of X ?

1-1	1-2	1-3	1-4	1-5	1-6
2-1	2-2	2-3	2-4	2-5	2-6
3-1	3-2	3-3	3-4	3-5	3-6
4-1	4-2	4-3	4-4	4-5	4-6
5-1	5-2	5-3	5-4	5-5	5-6
6-1	6-2	6-3	6-4	6-5	6-6

Example (Minimum of Three Dice) If X = the minimum of *three* fair six-sided dice, what is the mean value of X ? Find the answer by simulation in R.

In these examples, we were forced to calculate μ and σ from the definitions (or simulation). There are some random experiments that are so commonly used that statisticians have developed exact formulas for μ and σ that are easy to calculate. We turn to these next.

- Binomial
- Geometric
- Hypergeometric
- Poisson

4.2 - Binomial Distribution

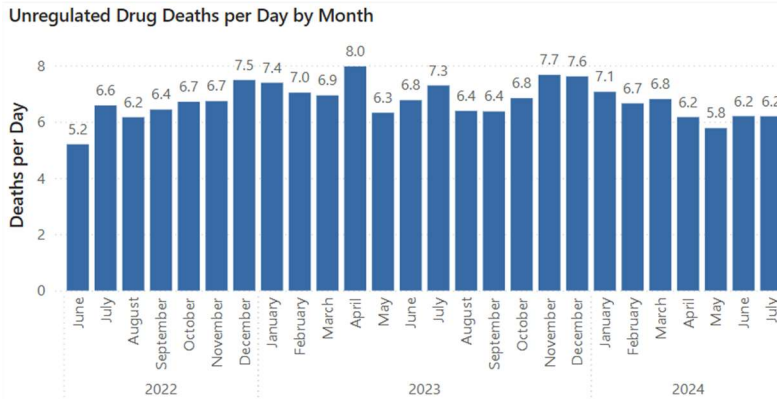
The binomial distribution arises when we are concerned with the variable:

X = the number of “successes” in a series of n independent trials

Here a *trial* (also called a *trial*) is any random experiment that has just two possible outcomes. By convention, these two outcomes are called “success” and “failure”. For instance, they could be:

- Win/Lose
- Live/Die
- True/False
- Pass/Fail
- Within Specification/Not Within Specification

Example (Drug Deaths) Suppose $n = 10\,000$ people consume opioid drugs today in BC. To simplify, suppose that each opioid user has the same risk $p = 0.06\%$ of dying today.



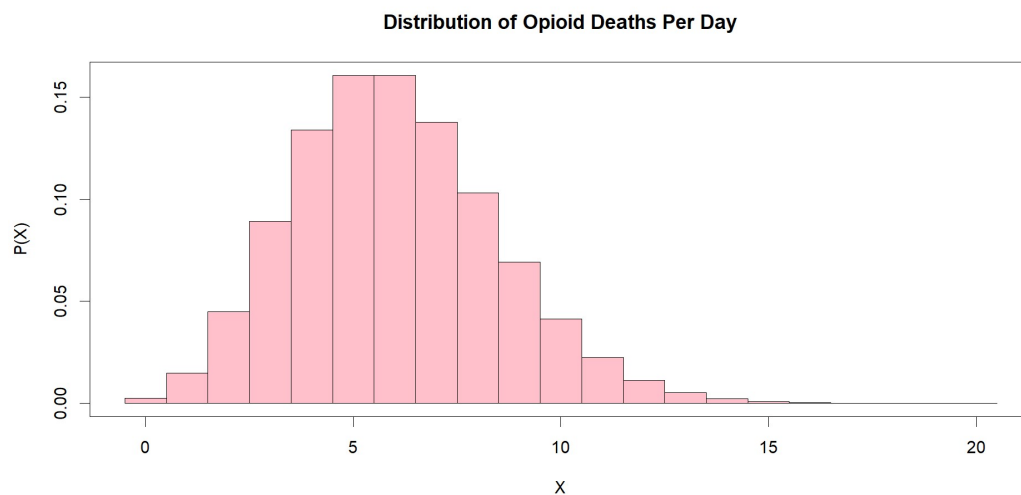
We make these assumptions:

1. **Fixed n** – the total number of users today is known and fixed.
2. **Success/Failure** – each opioid user either dies (“success”) or lives (“failure”).
3. **Equal p** – each opioid user has the same probability of dying today, $p = 0.0006$
4. **Independence** – each user’s outcome is unaffected by other users’ outcomes

Finally, let

X = the number of deaths today due to opioid overdose

(example continued) Under the above assumptions, the variable X follows a *binomial* distribution, shown here:



As an example, let's calculate just one value:

$$P(X = 5)$$

CAUTION: we need to verify that all of the conditions of the binomial distribution are satisfied before we apply our formulas in any given problem.

Example Rolling a fair die 5 times and counting the number of 3s obtained is an example of a binomial experiment. Here X = the number of 3s obtained.

Check the conditions:

Calculate the probability distribution of X .

General Formula If a variable X satisfies the conditions of a binomial variable, with n trials and probability p of success, then

$$P(x) = {}^nC_x \cdot p^x \cdot q^{n-x}$$

where

$${}^nC_x = \frac{n!}{(n-x)!x!}$$

Example Suppose 5 cards are selected *with* replacement from a deck. Find the probability that 2 are jacks.

Example Suppose 5 cards are selected *with* replacement from a deck. Find the probability that at least 2 are jacks.

Mean and Variance for Binomial Distributions

The mean and variance of any probability distribution are found using:

$$\mu = \sum_x [x \cdot P(x)]$$

$$\sigma^2 = \sum_x [(x - \mu)^2 \cdot P(x)]$$

It is possible to start with these definitions and to then use the binomial distribution assumptions to derive the results:

$$\mu = np$$

$$\sigma^2 = npq$$

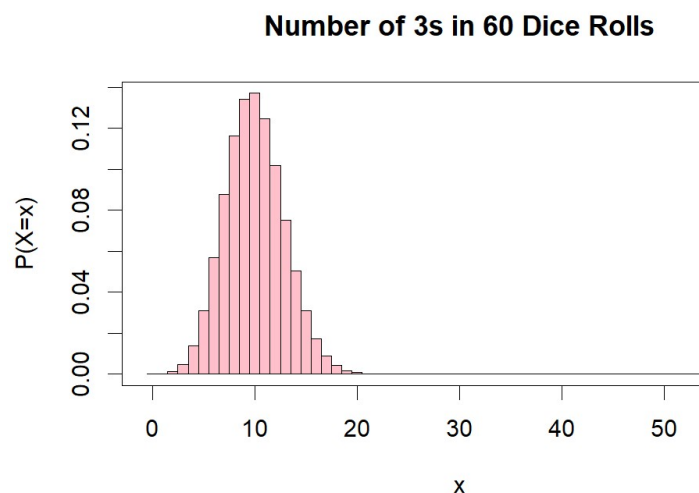
Example In the earlier example about opioid deaths, we had $n = 10\,000$ and $p = 0.0006$. Then the mean value of X (deaths in a day) is:

$$\mu = np = 10000 \times 0.0006 = 6$$

Example Suppose we roll a six-sided die 60 times. Let X = the number of times we roll a 3. Then the *parameters* of the distribution of X are:

$$\mu =$$

$$\sigma =$$



Example (Empirical Rule) For a certain model of laser printer, the probability of a unit needing repairs in the first year is 10%. In an attempt to reduce this rate, modifications were made to the printer design and a year later a study of 200 randomly selected modified laser printers was conducted.

- Assuming the modifications had no effect on the printer reliability, find the mean and standard deviation of the number of laser printers that need repair among 200.
- Suppose that in the sample of 200 laser printers, 14 needed repair. Assuming the modifications had no effect, is this rate unusually low?
- Does this sample provide good evidence that the modifications improved the reliability of this model of laser printer?

4.3 - Hypergeometric Distribution

The hypergeometric distribution is similar to the binomial distribution. In both cases, we are concerned with the variable

X = number of successes in a series of n trials

The difference is shown in the table below:

Binomial	Hypergeometric
<ul style="list-style-type: none"> • trials are <i>independent</i> • sampling <i>with replacement</i>, or sampling from an <i>infinite</i> population • same probability p for each trial 	<ul style="list-style-type: none"> • trials are <i>dependent</i> • sampling <i>without replacement</i> from a finite population • probability of success changes for subsequent trials

If the number of trials, n , is small compared to the population size, then the Binomial model and the Hypergeometric model give *very similar* probabilities.

Hypergeometric Distribution Formula

Suppose a population of N objects contains K “success” objects and $N - K$ “failure” objects. If you select a random sample of size n from this population, let

X = the number of “success” objects in the sample.

Then the probability of getting x “success” objects is

$$P(x) = \frac{C(K, x) \cdot C(N - K, n - x)}{C(N, n)}$$

for any $x = 0, 1, 2, \dots, n$.

The mean and variance of X are:

$$\mu = n \cdot \frac{K}{N}$$

$$\sigma^2 = n \cdot \frac{K}{N} \cdot \frac{N - K}{N} \cdot \frac{N - n}{N - 1}$$

Example Suppose you sample 5 cards from a deck of 52 cards. Let X = the number of Jacks you get in the sample? (There are 4 Jacks in the deck). What is $P(X = 2)$?

Hypergeometric (without replacement)

$$P(X = 2) = \frac{C(4, 2) \cdot C(48, 3)}{C(52, 5)} =$$

```
> dhyper(2, 4, 48, 5)
[1] 0.03992982
```

Binomial (with replacement)

$$P(X = 2) = C(5, 2) \cdot \left(\frac{4}{52}\right)^2 \cdot \left(\frac{48}{52}\right)^3 =$$

```
> dbinom(2, 5, 4/52)
[1] 0.04654006
```

Now, suppose you combine 10 decks of cards into one “superdeck”. If you randomly select 5 cards, what is the probability of getting 2 Jacks?

Hypergeometric:

Binomial:

5% Rule

If the sample size n is 5% or less of the total population N , then probabilities calculated using Hypergeometric and Binomial distributions are practically the same.

If $n > 0.05 \times N$ then be sure to use Hypergeometric.

Example Suppose a class of 90 students has 40 Apple Mac users. To conduct a market research survey, you randomly select a sample of 10 students.

Let X = the number of Apple Mac users in the sample.

Then X is a hypergeometric variable with parameters:

$$N =$$

$$K =$$


$$n =$$

- a. What is the probability that 4 students in the sample are Apple Mac users?

- b. What is the probability that 4 *or fewer* of them are Apple Mac users?

c. What are the mean and standard deviation of X ?

d. Generate a probability histogram of X .



e. If you found $X = 8$ Mac users, would this be considered *unusual*?

4.4 - Geometric Distribution

The Geometric Distribution is based on the same assumptions as the binomial distribution: n trials each with probability of success p . In this case, we are working with the variable

X = the number of trials it takes to obtain the first “success”

Example Suppose a tech-support telephone help line is occupied 75% of the time. Find the probability that you will have to call 3 times to gain access.

Solution: If you first gain access on the 3th call, then you had to have:

- failure on the first call: probability $q = 0.75$
- failure on the second call: probability $q = 0.75$
- success on the third call: probability $p = 0.25$

Therefore,

$$P(3) =$$

Geometric Distribution Formulas

Suppose every trial of a random experiment has probability p of success and probability $q = 1 - p$ of failure, and all trials are independent.

Let X = the number of trials it takes to get the first success (possible values: 1, 2, 3, ...). The probability distribution is given by:

$$P(X = x) = q^{x-1} \cdot p$$

The mean value and variance of X are consequently:

$$\mu = \frac{1}{p}$$
$$\sigma^2 = \frac{1-p}{p^2}$$

Derivation of μ

You are not expected to be able to perform a derivation like the following, but it may help you to be familiar with it.

$\mu = \sum_{x=1}^{\infty} x \cdot P(x)$	definition of mean value
$= \sum_{x=1}^{\infty} x \cdot q^{x-1} \cdot p$	sub formula for $P(x)$
$= p \cdot \left[\sum_{x=1}^{\infty} x \cdot q^{x-1} \right]$	factor out the p
$= p \cdot \frac{d}{dq} \left[\sum_{x=1}^{\infty} q^x \right]$	since $x \cdot q^{x-1} = \frac{d}{dq} [q^x]$
$= p \cdot \frac{d}{dq} \left[\frac{q}{1-q} \right]$	sum of a geometric series
$= p \cdot \frac{(1-q) + q}{(1-q)^2}$	Quotient Rule
$= p \cdot \frac{1}{(1-q)^2}$	
$= \frac{p}{p^2} = \frac{1}{p}$	algebraic simplification (using $p = 1 - q$)

Example Suppose a tech-support telephone help line is occupied 75% of the time. Let X = the number of times you must call until you get access.

Find $E[X] =$

Find $\sigma_X =$

(example continued)

Find $P(X \geq 4)$

Generate a probability histogram for the variable X using R.

4.5 - Poisson Distribution

The *Poisson Distribution* is named after the French mathematician Siméon Poisson (1781-1840). This distribution is one of the most important probability distributions in engineering and computer science. It arises whenever we are modelling events that occur randomly throughout a given time interval.

Examples The Poisson distribution could be applied for variables X like the following:

X = number of jobs arriving for service at a CPU per second

X = number of bits arriving in error at a network node per minute

Assumptions for Poisson Distribution

- X = the number of occurrences of an event over some interval.
- Occurrences happen with uniform probability over the time interval.
- Occurrences are independent of each other.
- The *mean* number of occurrences is known:

λ = mean number of occurrences during the time interval

Poisson Distribution Formulas

Given the assumption above, the probability of having x occurrences is:

$$P(X = x) = \frac{\lambda^x}{x!} e^{-\lambda}$$

Here $e = 2.71828 \dots$ is Euler's constant from calculus. The mean and variance of X are:

$$\mu = \lambda$$

$$\sigma^2 = \lambda$$

Example Suppose a web server gets on average one request every 2 seconds. Assuming requests arrive randomly and independently over time, what is the probability that 20 requests will arrive during a one-minute interval?

What is the probability that 20 *or fewer* requests will arrive in a given minute?

Is getting 20 requests in a given minute *unusual* in the statistical sense?

(example continue) Generate the probability histogram for X = the number of web requests in a given minute.