

Processes and Threads

Chapter 2

The Process Model

A process is just an instance of an executing program, including the current values of the program counter (PC), registers, and variables.

The Process Model

Example:

Code

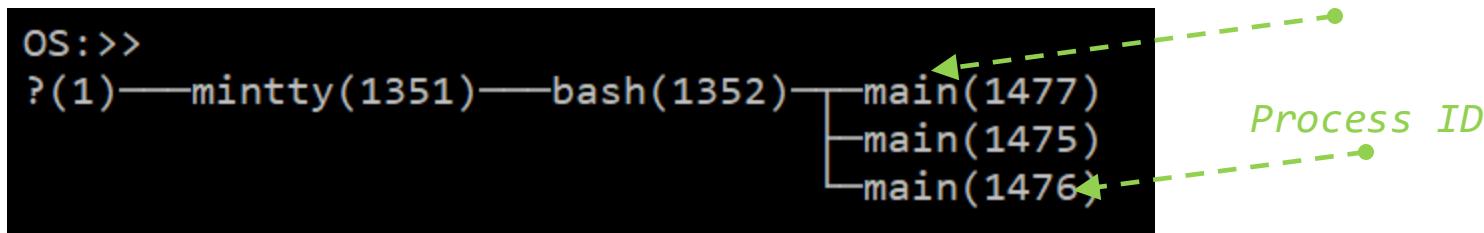
```
#include <stdio.h>
#include <unistd.h>

int main()
{
    printf(" Preocess ID : %ld \n", getpid());
    sleep(100);
    return 0;
}
```

Location

Name	Date modified	Type	Size
main.c	2020-07-19 3:09 PM	C File	1 KB
main.exe	2020-07-19 3:06 PM	Application	164 KB

Processes



The Process Model

```
#include <stdio.h>
int main()
{
    int sum = 0;
    for(int i=0; i<5; i++){
        sum += i;
    }
    return 0;
}
```

The Process Model

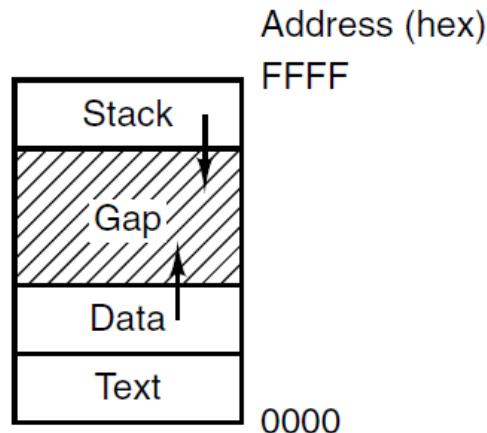
```
#include <stdio.h>
int main()
{
    int sum = 0;
    for(int i=0; i<5; i++){
        sum += i;
    }
    return 0;
}
```

Memory Location

Assembly code	
401102	main: 55 push rbp
401103	48 89 e5 mov rbp,rsp
401106	c7 45 fc 00 00 00 00 mov DWORD PTR [rbp-0x4],0x0
40110d	c7 45 f8 00 00 00 00 mov DWORD PTR [rbp-0x8],0x0
401114	83 7d f8 04 cmp DWORD PTR [rbp-0x8],0x4
401118	7f 0c jg 401126 <main+0x24>
40111a	8b 45 f8 mov eax,DWORD PTR [rbp-0x8]
40111d	01 45 fc add DWORD PTR [rbp-0x4],eax
401120	83 45 f8 01 add DWORD PTR [rbp-0x8],0x1
401124	eb ee jmp 401114 <main+0x12>
401126	b8 00 00 00 00 mov eax,0x0
40112b	5d pop rbp
40112c	c3 ret
40112d	0f 1f 00 nop DWORD PTR [rax]

The Process Model

Processes in UNIX have their memory divided up into three segments: the text segment (i.e., the program code), the data segment (i.e., the variables), and the stack segment. The data segment grows upward and the stack grows downward.



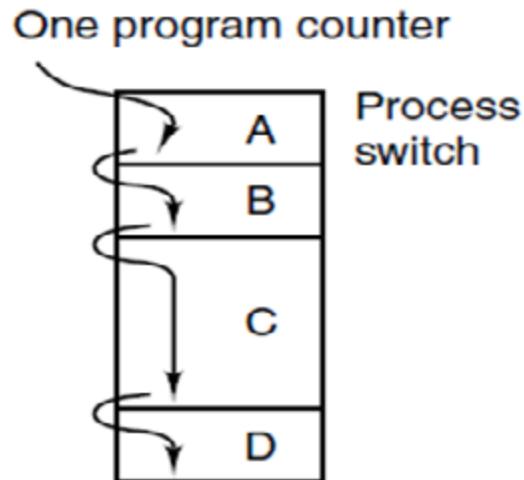
Processes have three segments: text, data, and stack.

The Process Model

Multiprogramming system

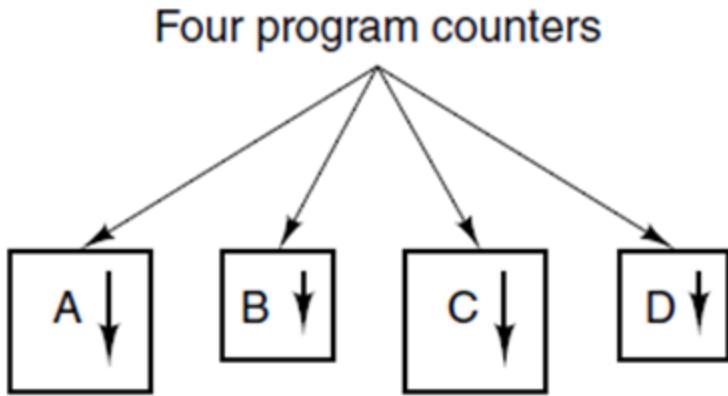
1. The CPU switches from process to process quickly.
2. The CPU runs each process for tens or hundreds of milliseconds.
3. Strictly speaking, at any one instant the CPU is running only one process.

The rapid switching back and forth is called multiprogramming.



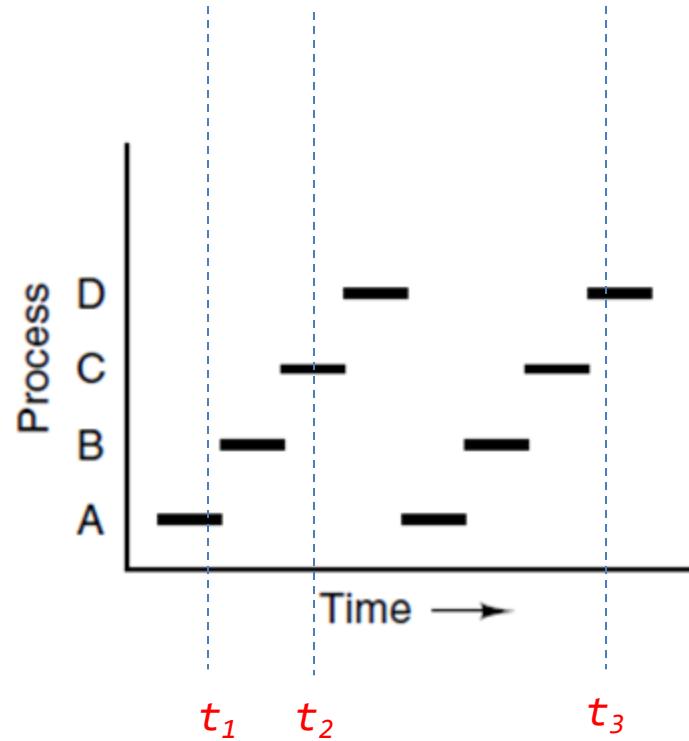
Multiprogramming of four programs

The Process Model



Conceptual model of
four independent, sequential processes.

The Process Model



Only one program is active at once.

Process Creation

Four principal events that cause processes to be created:

1. System initialization.
2. Execution of a process-creation system call by a running process.
3. A user request to create a new process.
4. Initiation of a batch job.

Process Creation

System initialization

When an operating system is booted, typically numerous processes are created.

- 1) Some of these are foreground processes, interacting with (human) users and perform work for them.
- 2) Others run in the background and are not associated with particular users, but instead have some specific function.

Processes that stay in the background to handle some activity such as email, Web pages, news, printing, and so on are called daemons.

Task Manager

Windows 10

- □ ×

File Options View

Processes Performance App history Startup Users Details Services

Name	PID	Description	Status	Group
vpnagent	29180	Cisco AnyConnect Secure Mobility Agent	Running	
WpnUserService_1c2879	6508	Windows Push Notifications User Service_1c2879	Running	UnistackSvcGroup
UserDataSvc_1c2879	5224	User Data Access_1c2879	Running	UnistackSvcGroup
UnistoreSvc_1c2879	5224	User Data Storage_1c2879	Running	UnistackSvcGroup
PrintWorkflowUserSvc_1c2...	30772	PrintWorkflow_1c2879	Running	PrintWorkflow
PimlIndexMaintenanceSvc_...	5224	Contact Data_1c2879	Running	UnistackSvcGroup
OneSyncSvc_1c2879	5224	Sync Host_1c2879	Running	UnistackSvcGroup
CDPUserSvc_1c2879	6416	Connected Devices Platform User Service_1c2879	Running	UnistackSvcGroup
BcastDVRUserService_1c2879	28176	GameDVR and Broadcast User Service_1c2879	Running	BcastDVRUserService
wuauserv	6160	Windows Update	Running	netsvcs
WSearch	9120	Windows Search	Running	
wscsvc	1560	Security Center	Running	LocalServiceNetworkRestricted
WpnService	3896	Windows Push Notifications System Service	Running	netsvcs
wlidsvc	1792	Microsoft Account Sign-in Assistant	Running	netsvcs
WlanSvc	2972	WLAN AutoConfig	Running	LocalSystemNetworkRestricted
Winmgmt	2416	Windows Management Instrumentation	Running	netsvcs
WinHttpAutoProxySvc	41416	WinHTTP Web Proxy Auto-Discovery Service	Running	LocalServiceNetworkRestricted
WinDefend	3884	Windows Defender Antivirus Service	Running	
WdNisSvc	5164	Windows Defender Antivirus Network Inspection Service	Running	
WdiSystemHost	9268	Diagnostic System Host	Running	LocalSystemNetworkRestricted
WdiServiceHost	4724	Diagnostic Service Host	Running	LocalService
wcncsvc	7668	Windows Connect Now - Config Registrar	Running	LocalServiceAndNoImpersonation
Wcmsvc	2788	Windows Connection Manager	Running	LocalServiceNetworkRestricted
WbioSrvc	16936	Windows Biometric Service	Running	WbioSvcGroup
W32Time	1208	Windows Time	Running	LocalService
VaultSvc	768	Credential Manager	Running	
UsoSvc	6160	Update Orchestrator Service	Running	netsvcs
UserManager	2120	User Manager	Running	netsvcs
TrustedInstaller	16684	Windows Modules Installer	Running	
TrkWks	3872	Distributed Link Tracking Client	Running	LocalSystemNetworkRestricted
TPHKLOAD	3848	Lenovo Hotkey Client Loader	Running	
TokenBroker	7796	Web Account Manager	Running	netsvcs
TimeBrokerSvc	1324	Time Broker	Running	LocalServiceNetworkRestricted
Themes	1936	Themes	Running	netsvcs

Fewer details | Open Services

Ubuntu 18.04.4 LTS

ps -aux | less

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.1	0.1	225984	9744	?	Ss	14:31	0:11	/sbin/init splash
root	2	0.0	0.0	0	0	?	S	14:31	0:00	[kthreadd]
root	6	0.0	0.0	0	0	?	I<	14:31	0:00	[mm_percpu_wq]
root	7	0.0	0.0	0	0	?	S	14:31	0:00	[ksoftirqd/0]
root	8	0.4	0.0	0	0	?	I	14:31	0:33	[rcu_sched]
root	9	0.0	0.0	0	0	?	I	14:31	0:00	[rcu_bh]
root	10	0.0	0.0	0	0	?	S	14:31	0:00	[migration/0]
root	11	0.0	0.0	0	0	?	S	14:31	0:00	[watchdog/0]
root	12	0.0	0.0	0	0	?	S	14:31	0:00	[cpuhp/0]
root	13	0.0	0.0	0	0	?	S	14:31	0:00	[cpuhp/1]
root	14	0.0	0.0	0	0	?	S	14:31	0:00	[watchdog/1]
root	15	0.0	0.0	0	0	?	S	14:31	0:00	[migration/1]
root	16	0.0	0.0	0	0	?	S	14:31	0:03	[ksoftirqd/1]
root	19	0.0	0.0	0	0	?	S	14:31	0:00	[cpuhp/2]
root	20	0.0	0.0	0	0	?	S	14:31	0:00	[watchdog/2]
root	21	0.0	0.0	0	0	?	S	14:31	0:00	[migration/2]
root	22	0.0	0.0	0	0	?	S	14:31	0:00	[ksoftirqd/2]
root	24	0.0	0.0	0	0	?	I<	14:31	0:00	[kworker/2:0H]
root	25	0.0	0.0	0	0	?	S	14:31	0:00	[cpuhp/3]
root	26	0.0	0.0	0	0	?	S	14:31	0:00	[watchdog/3]
root	27	0.0	0.0	0	0	?	S	14:31	0:00	[migration/3]
root	28	0.0	0.0	0	0	?	S	14:31	0:00	[ksoftirqd/3]
root	30	0.0	0.0	0	0	?	I<	14:31	0:00	[kworker/3:0H]
root	31	0.0	0.0	0	0	?	S	14:31	0:00	[kdevtmpfs]
root	32	0.0	0.0	0	0	?	I<	14:31	0:00	[netns]
root	33	0.0	0.0	0	0	?	S	14:31	0:00	[rcu_tasks_kthre]
root	34	0.0	0.0	0	0	?	S	14:31	0:00	[kauditfd]
root	39	0.0	0.0	0	0	?	S	14:31	0:00	[khungtaskd]
root	40	0.0	0.0	0	0	?	S	14:31	0:00	[oom_reaper]
root	41	0.0	0.0	0	0	?	I<	14:31	0:00	[writeback]
root	42	0.0	0.0	0	0	?	S	14:31	0:00	[kcompactd0]
root	43	0.0	0.0	0	0	?	SN	14:31	0:00	[ksmd]
root	44	0.0	0.0	0	0	?	SN	14:31	0:00	[khugepaged]
root	45	0.0	0.0	0	0	?	I<	14:31	0:00	[crypto]

Process Creation

Execution of a process creation system call
by a running process

Often a running process will issue system calls to create one or more new processes to help it do its job.

For example, if a large amount of data is being fetched over a network for subsequent processing, it may be convenient to create one process to fetch the data and put them in a shared buffer while a second process removes the data items and processes them. On a multiprocessor, allowing each process to run on a different CPU may also make the job go faster.

Process Creation

A user request to create a new process

In interactive systems, users can start a program by typing a command or (double) clicking on an icon. Taking either of these actions starts a new process and runs the selected program in it.

In command-based UNIX systems running X, the new process takes over the window in which it was started.

```
OS:>> ls
main.c main.exe
OS:>> ./main.exe
    Preocess ID : 1487
OS:>>
```

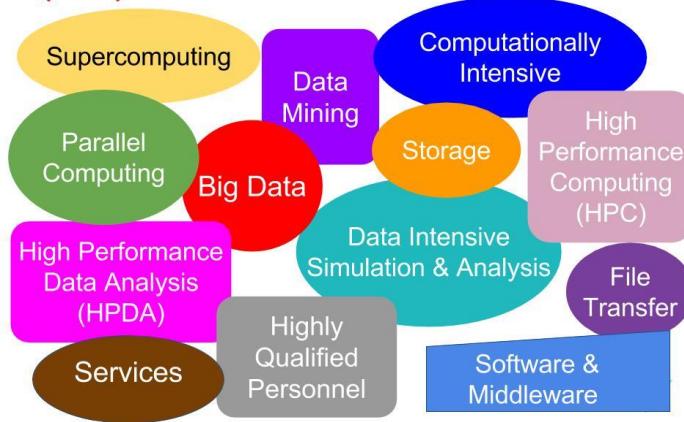
Process Creation

Initiation of a batch job

The last situation in which processes are created applies only to the batch systems found on large mainframes. Think of inventory management at the end of a day at a chain of stores. Here users can submit batch jobs to the system (possibly remotely).

When the operating system decides that it has the resources to run another job, it creates a new process and runs the next job from the input queue in it.

What is Advanced Research Computing (ARC)?



Digital Research Alliance of Canada | Alliance de recherche numérique du Canada

About ▾ Membership ▾ Services ▾ Funding Opportunities Initiatives ▾ Lat

Strategic Plan 2022-2025

The Alliance's inaugural strategic plan is a commitment to improving access to digital tools and services for all Canadian researchers.

This diagram shows a complex network of interconnected nodes representing various research tools and data analysis methods. Nodes include a smartphone, a laptop displaying a bar chart, a server rack, a cloud icon, a microscope, a telescope, a database, a monitor showing a waveform, and a small screen displaying binary code. Lines connect these nodes to represent the flow of data and collaboration between different research domains.

<https://alliancecan.ca/en>

Process Termination

Typical conditions which terminate a process:

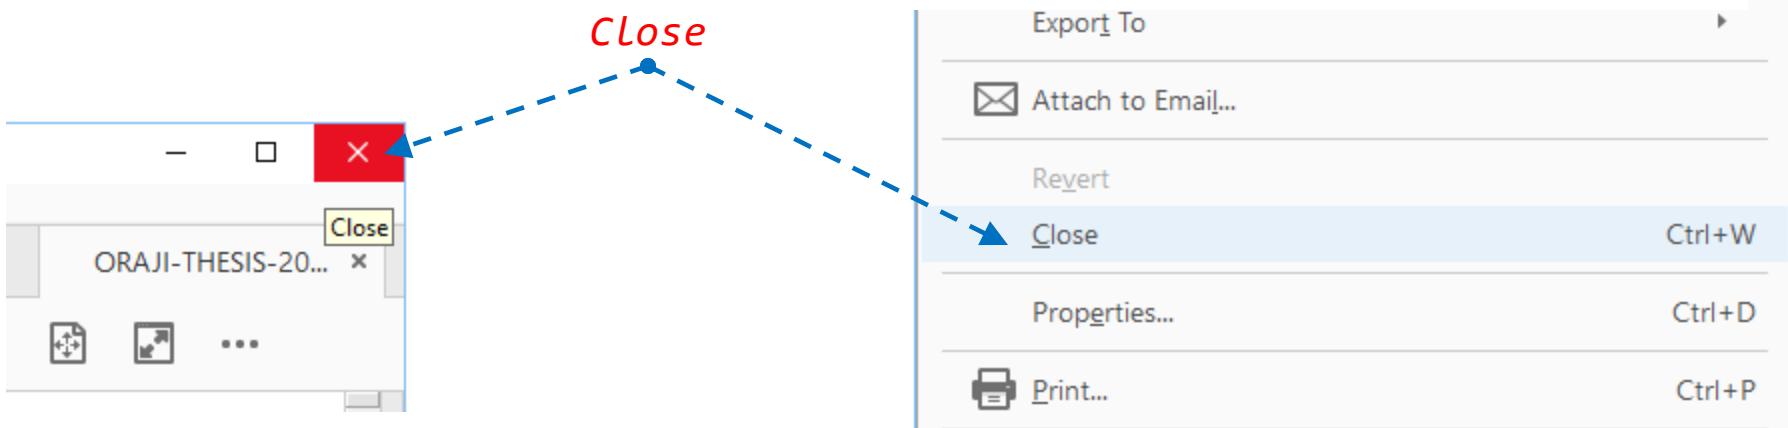
1. Normal exit (voluntary).
2. Error exit (voluntary).
3. Fatal error (involuntary).
4. Killed by another process (involuntary).

Process Termination

Normal exit (voluntary)

Most processes terminate because they have done their work. When a compiler has compiled the program given to it, the compiler executes a system call to tell the operating system that it is finished. This call is exit in UNIX and ExitProcess in Windows.

Screen-oriented programs also support voluntary termination. Word processors, Internet browsers, and similar programs always have an icon or menu item that the user can click to tell the process to remove any temporary files it has open and then terminate.



Process Termination

Error exit (voluntary)

The second reason for termination is that the process discovers a fatal error. For example, if a user types the command `cc foo.c`

to compile the program `foo.c` and no such file exists, the compiler simply announces this fact and exits.

```
~$ gcc foo.c
gcc: error: foo.c: No such file or directory
gcc: fatal error: no input files
compilation terminated.
~$
```

Screen-oriented interactive processes generally do not exit when given bad parameters. Instead they pop up a dialog box and ask the user to try again.

Process Termination

Fatal error (involuntary)

The third reason for termination is an error caused by the process, often due to a program bug. Examples include executing an illegal instruction, referencing nonexistent memory, or dividing by zero.

In some systems (e.g., UNIX), a process can tell the operating system that it wishes to handle certain errors itself, in which case the process is signaled (interrupted) instead of terminated when one of the errors occurs.

Process Termination

Killed by another process (involuntary)

The fourth reason a process might terminate is that the process executes a system call telling the operating system to kill some other process. In UNIX this call is kill. The corresponding Win32 function is TerminateProcess. In both cases, the killer must have the necessary authorization to do in the killee.

```
OS:>> ./program.exe &
[1] 1243
OS:>> Process waiting 30 seconds.

OS:>> pstree -p
?(1)---mintty(1169)---bash(1170)---program(1243)
                                         \pstree(1244)

OS:>> kill 1243
OS:>> fg
-bash: fg: job has terminated
[1]+  Terminated                  ./program.exe
OS:>>
```

Process Hierarchies

In some systems, when a process creates another process, the parent process and child process continue to be associated in certain ways. The child process can itself create more processes.

Process Hierarchies

UNIX

In UNIX, a process and all of its children and further descendants together form a process group.

When a user sends a signal from the keyboard, the signal is delivered to all members of the process group currently associated with the keyboard.

Individually, each process can catch the signal, ignore the signal, or take the default action, which is to be killed by the signal.

Process Hierarchies

Windows

In contrast, Windows has no concept of a process hierarchy. All processes are equal.

The only hint of a process hierarchy is that when a process is created, the parent is given a special token (called a **handle**) that it can use to control the child.

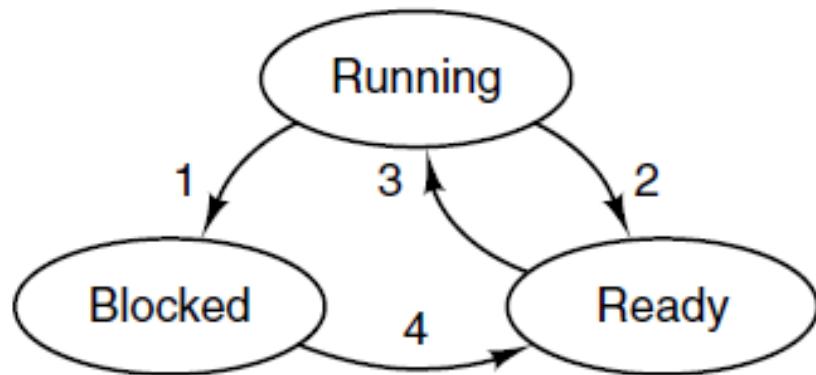
However, it is free to pass this token to some other process, thus invalidating the hierarchy. Processes in UNIX cannot disinherit their children.

Process States

Three states a process may be in:

1. Running (actually using the CPU at that instant)
2. Ready (scheduled and runnable; temporarily stopped to let another process run)
3. Blocked (unable to run until some external event happens)

Process States



1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available

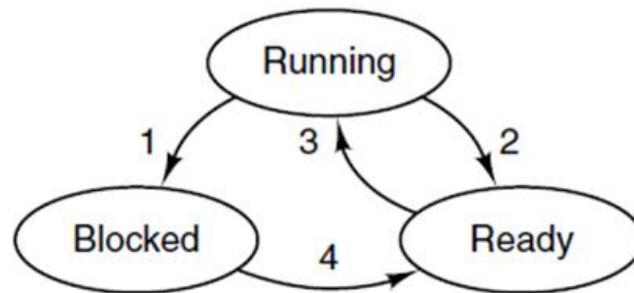
A process can be in running, blocked, or ready state. Transitions between these states are as shown.

Process States

Transition 1 occurs when the operating system discovers that a process cannot continue right now.

In some systems the process can execute a system call, such as pause, to get into blocked state.

In other systems, including UNIX, when a process reads from a pipe or special file (e.g., a terminal) and there is no input available, the process is automatically blocked.



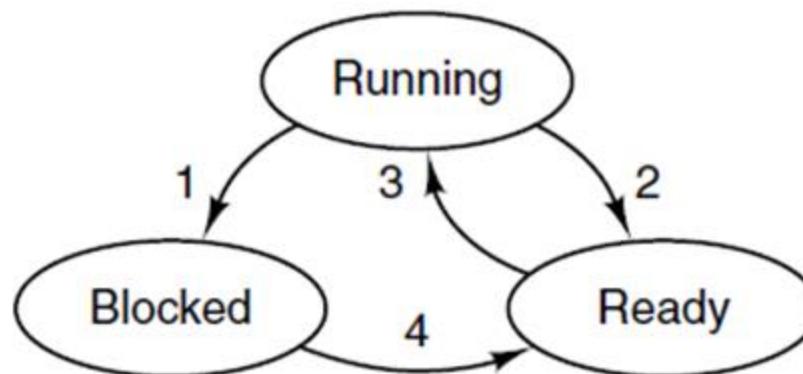
```
OS:>>
OS:>> cat data.txt | tr '[A-Z]' 'a-z'| tr -C '[a-z]' '\n' | grep cow
cow          1
cows         2
cows         3
cows         4
```

Process States

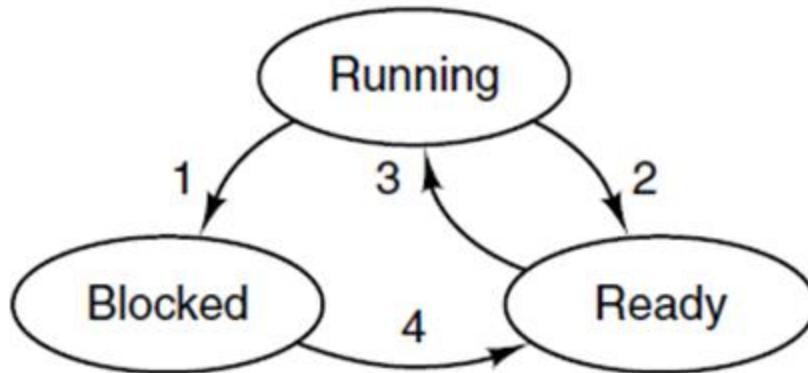
Transitions 2 and 3 are caused by the process scheduler, a part of the operating system, without the process even knowing about them.

Transition 2 occurs when the scheduler decides that the running process has run long enough, and it is time to let another process have some CPU time.

Transition 3 occurs when all the other processes have had their fair share and it is time for the first process to get the CPU to run again.



Process States



Transition 4 occurs when the external event for which a process was waiting (such as the arrival of some input) happens.

If no other process is running at that instant, transition 3 will be triggered and the process will start running.

Otherwise it may have to wait in *ready* state for a little while until the CPU is available and its turn comes.

Implementation of Processes

To implement the process model, the operating system maintains a table (an array of structures), called the **process table**, with one entry per process.

This entry contains important information about the process' state, including its program counter, stack pointer, memory allocation, the status of its open files, its accounting and scheduling information, and everything else about the process that must be saved when the process is switched from running to ready or blocked state so that it can be restarted later as if it had never been stopped.

Implementation of Processes

Context Switch

A **context switch** refers to the process of saving the current state of a running process and loading the saved state of another process so that it can resume execution.

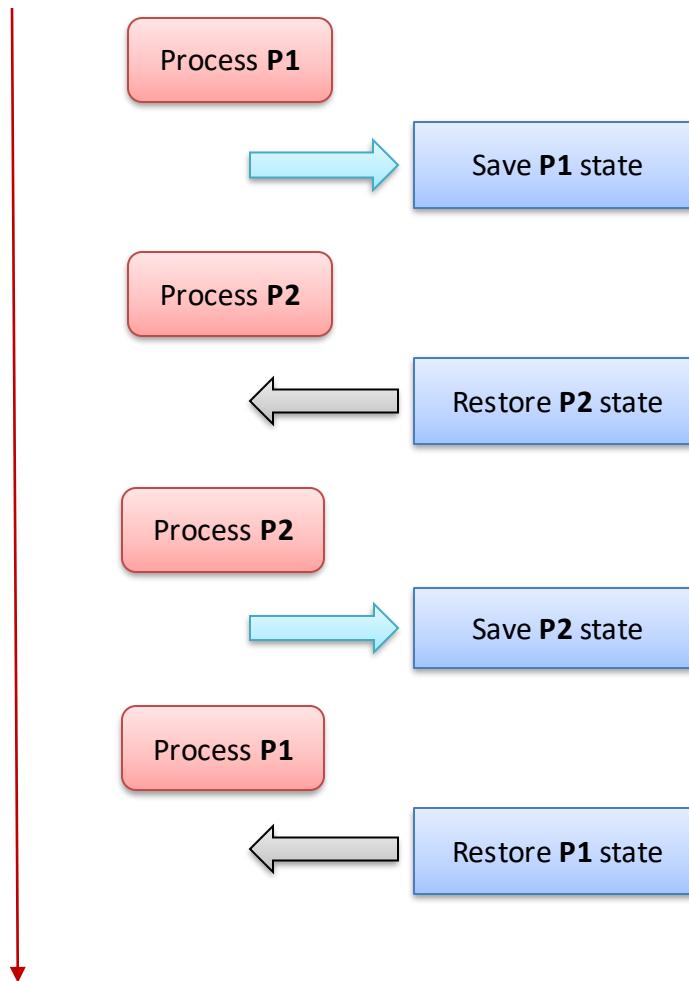
When the CPU scheduler decides to switch from one process to another, it must save the current process's context, including its program counter, register values, and other relevant state information, into the process's **PCB (Process Control Block)**.

PCB is a data structure that represents a single process in the operating system. The PCB is a crucial component within the process table.

Implementation of Processes

Context Switch

CPU



Implementation of Processes

Context Switch

```
1 #include <stdio.h>
2 #include <ucontext.h>
3 #include <unistd.h>
4 #define REG_RIP 16
5
6 /* user-level context (uc), includes
7 1) signal mask,
8 2) stack info,
9 3) link to next context,
10 and an embedded machine-level context (mc).
11
12 Machine-level context (mc) contains
13 General-purpose registers: RIP, RSP, RAX, etc.
14 . . . .
15 */
16 int main(int argc, const char *argv[]){
17     ucontext_t ctx;
18     printf("Capturing context...\n");
19     getcontext(&ctx);
20     printf("Context captured. Ready to inspect.\n");
21     printf("RIP: %llx\n", (unsigned long long)ctx.uc_mcontext.gregs[REG_RIP]);
22     return 0;
23 }
```

```
Capturing context...
Context captured. Ready to inspect.
RIP: 56f2d961c1f2
```

Implementation of Processes

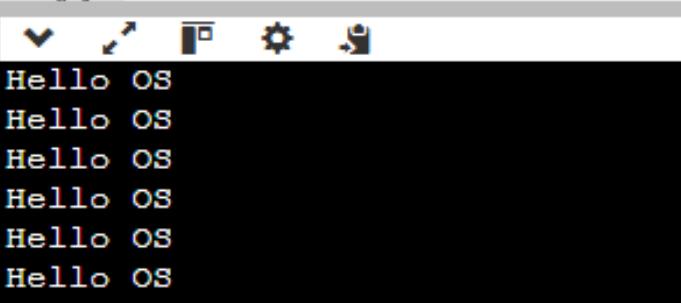
Context Switch

```
▲ start ■ pause ▶ continue ▷ step over ▷ step into ▷ step out ⓘ help
p ctx
$5 = {uc_flags = 140737353863488, uc_link = 0xffffffff, uc_stack = {
    ss_sp = 0x7fffff7daec40, ss_flags = -6128, ss_size = 0}, uc_mcontext = {
        gregs = {140737488349184, 140737488349192, 140733193388039,
            140737488349264, 140737353899408, 0, 140737488349192, 140737353899408,
            140737353866368, 140737354131072, 140737353895936, 140737488349200,
            140737353899034, 140737354117984, 140737353944517, 140737353897360,
            140737354083125, 676612543333817696, 140737352283740, 1,
            -3866529821160794880, 0, 1}, fpregs = 0x7fffffff880, __reserved1 = {
                140737353225526, 8388608, 140737352255748, 17179863168,
                14580214252548756736, 73325480322620, 140737353991104,
                140737488350016}}, uc_sigmask = {__val = {140737354037399, 0, 0, 0, 1,
                    140737351620416, 140737354131672, 140737354128048, 140737354088359, 13,
                    140737354130144, 73325480293968, 3, 17179869184, 2, 1241245548544}}},
    __fpregs_mem = { cwd = 289, swd = 0, ftw = 0, fop = 0, rip = 140737353888296,
        rdp = 3573412790848, mxcsr = 1, mxcr_mask = 0, _st = {{significand = {0,
            0, 0, 0}, exponent = 0, __glibc_reserved1 = {0, 0, 0}}, {
                significand = {2432, 0, 2432, 0}, exponent = 2432,
                __glibc_reserved1 = {0, 2432, 0}}, {significand = {2432, 0, 2432, 0},
                    exponent = 2432, __glibc_reserved1 = {0, 2432, 0}}, {significand = {
                        2432, 0, 2432, 0}, exponent = 2432, __glibc_reserved1 = {0, 2432,
                            0}}, {significand = {2432, 0, 2432, 0}, exponent = 2432,
                                __glibc_reserved1 = {0, 2432, 0}}, {significand = {
                                    2432, 0, 2432, 0}, exponent = 2432, __glibc_reserved1 = {0, 2432,
                                        0}}, {significand = {2432, 0, 2432, 0}, exponent = 2432,
                                            __glibc_reserved1 = {0, 2432, 0}}, {significand = {
                                                2432, 0, 2432, 0}, exponent = 2432, __glibc_reserved1 = {0, 2432,
                                                    0}}}, {significand = {2432, 0, 2432, 0}, exponent = 2432,
                                                        __glibc_reserved1 = {0, 2432, 0}}, {significand = {2432, 0, 2432, 0},
                                                            exponent = 2432, __glibc_reserved1 = {0, 2432, 0}}, {significand = {
                                                                2432, 0, 2432, 0}, exponent = 2432, __glibc_reserved1 = {0, 2432,
                                                                    0}}, {significand = {2432, 0, 2432, 0}, exponent = 2432,
                                                                        __glibc_reserved1 = {0, 2432, 0}}, {significand = {
                                                                            2432, 0, 2432, 0}, exponent = 2432, __glibc_reserved1 = {0, 2432,
                                                                                0}}}, {significand = {256, 0, 0, 0}, exponent = 0, __glibc_reserved1 = {0,
                                    64, 0}}}, xmm = {{element = {512, 1024, 0, 0}}, {element = {0, 0,
```

Implementation of Processes

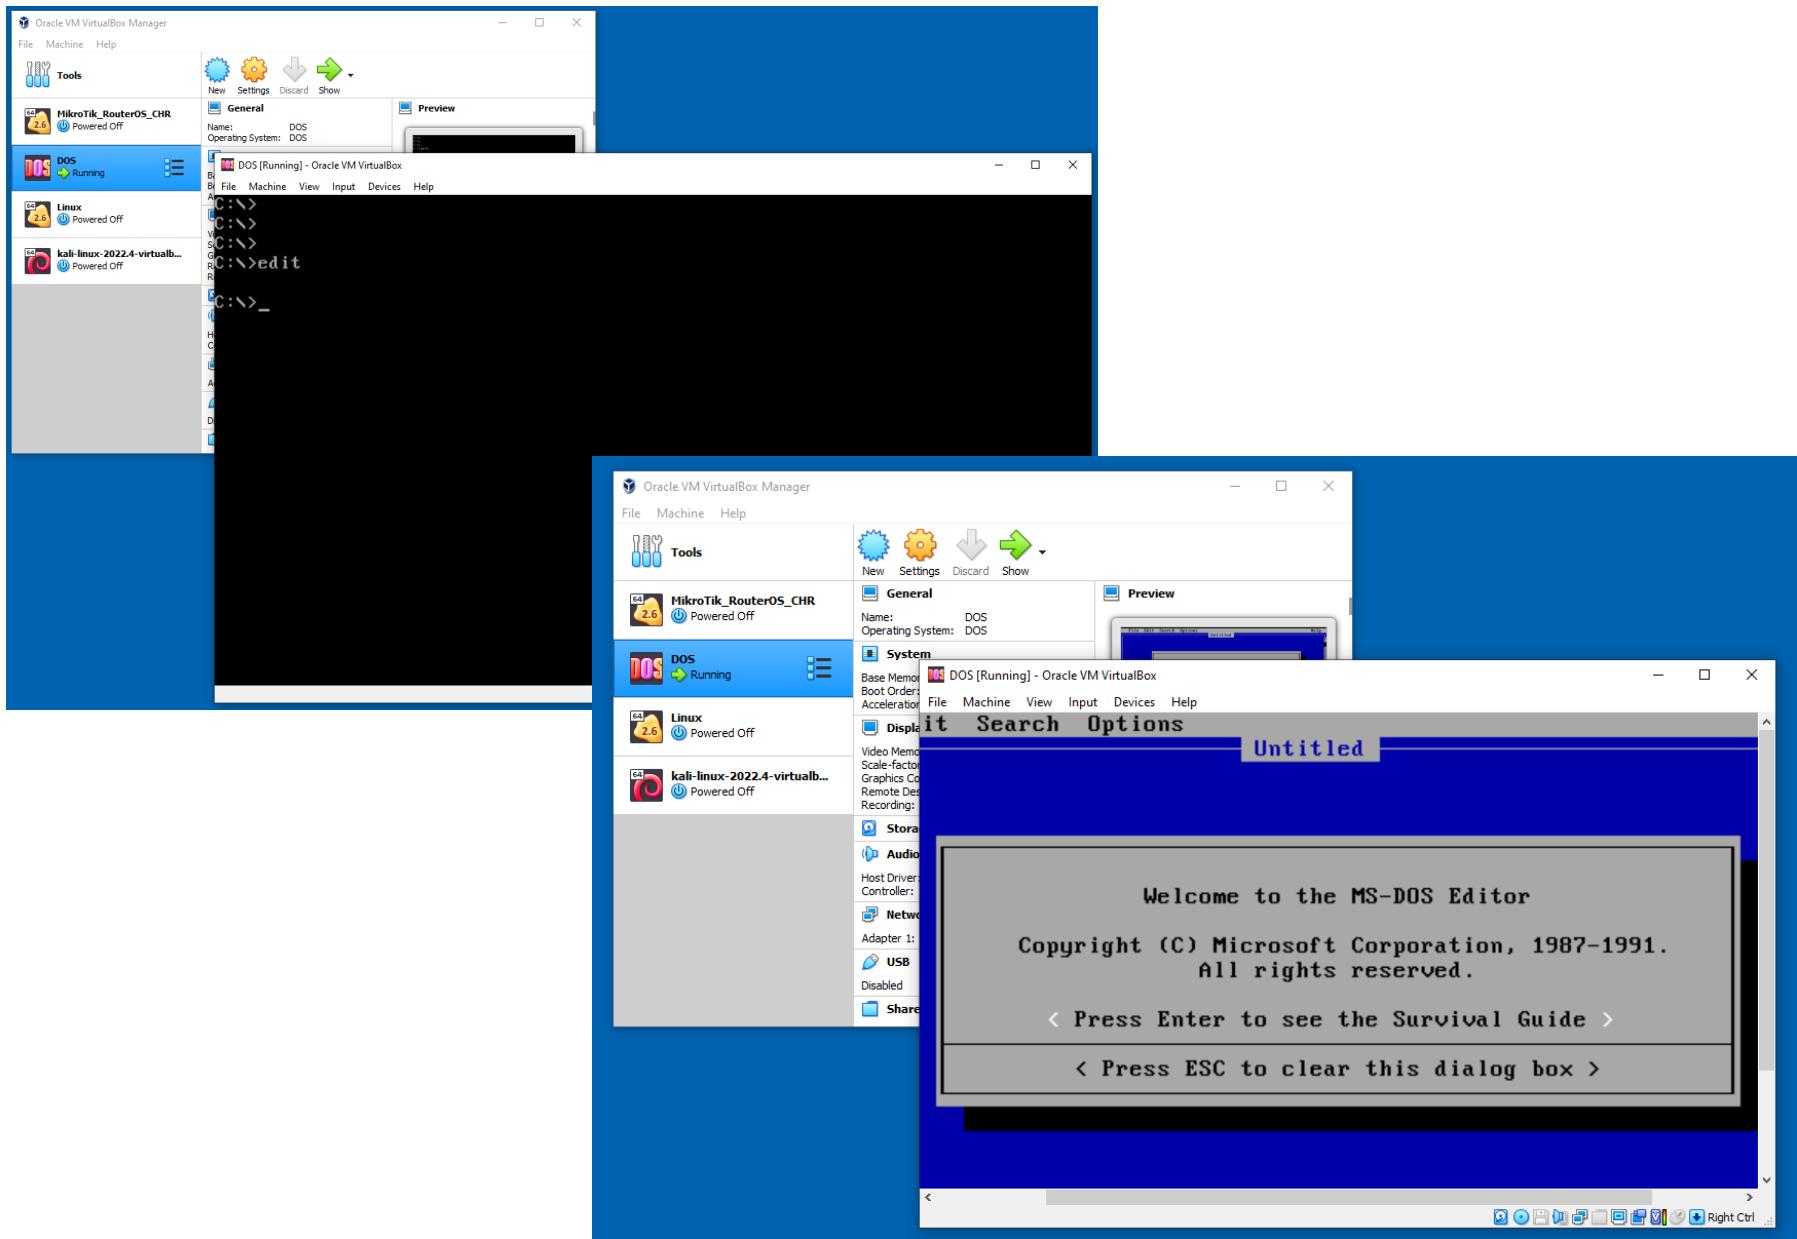
Context Switch

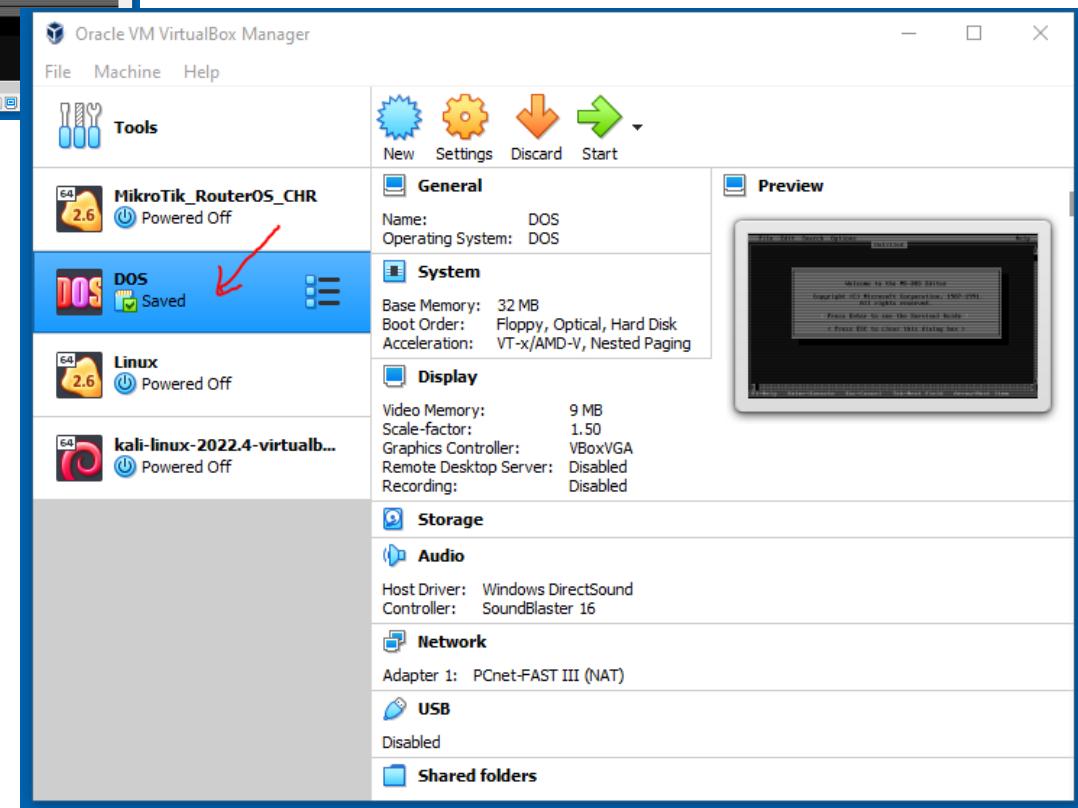
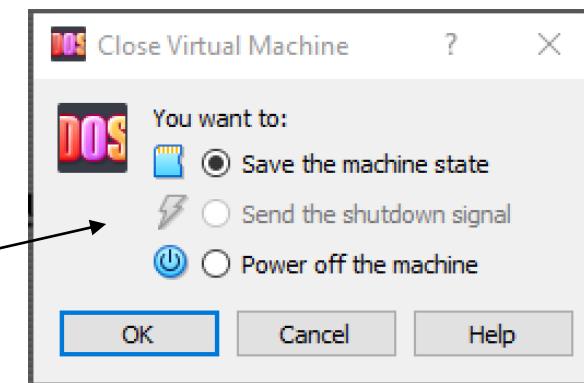
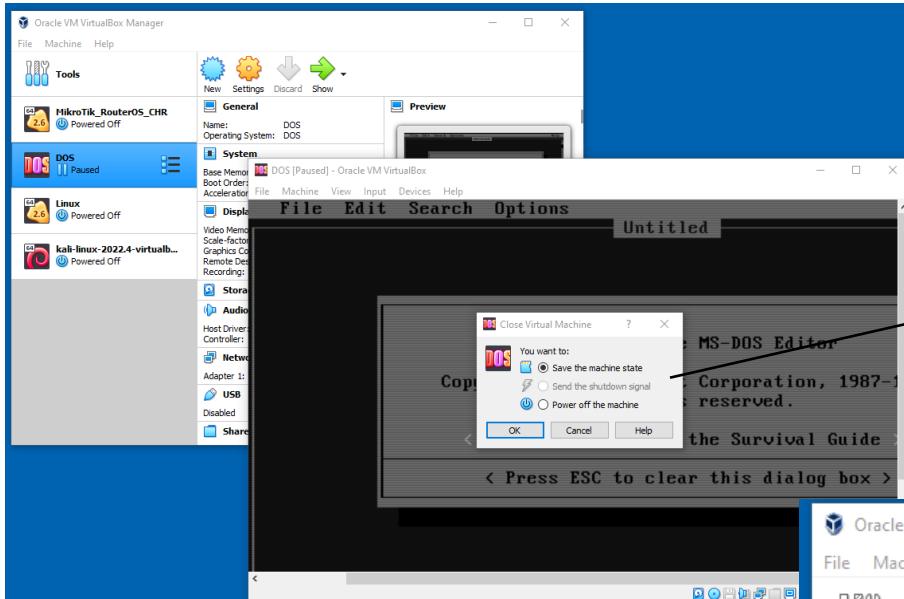
```
1 #include <stdio.h>
2 #include <ucontext.h>
3 #include <unistd.h>
4
5 int main(){
6     ucontext_t context;
7
8     getcontext(&context);
9     puts("Hello OS ");
10    sleep(1);
11    setcontext(&context);
12    return 0;
13 }
```



The screenshot shows a terminal window with a black background and white text. At the top, there is a toolbar with several icons: a downward arrow, a left arrow, a right arrow, a square, a gear, and a magnifying glass. Below the toolbar, the text "Hello OS" is repeated six times, each on a new line. The terminal window has a light gray border.

```
Hello OS
Hello OS
Hello OS
Hello OS
Hello OS
Hello OS
```





Oracle VM VirtualBox Manager

File Machine Help



Tools

New Settings Discard Start

MikroTik_RouterOS_CHR
2.6 Powered Off

DOS DOS Saved

Linux 2.6 Powered Off

kali-linux-2022.4-virtualb...
2.6 Powered Off

General
Name: DOS
Operating System: DOS

System
Base Memory: 32 MB
Boot Order: Floppy, Optical, Hard Disk
Acceleration: VT-x/AMD-V, Nested Paging

Display
Video Memory: 9 MB
Scale-factor: 1.50
Graphics Controller: VBoxVGA
Remote Desktop Server: Disabled
Recording: Disabled

Storage

Audio
Host Driver: Windows DirectSound
Controller: SoundBlaster 16

Network

Adapter 1: PCnet-FAST III (NAT)

USB

Disabled

Shared folders

Preview
 Welcome to the MS-DOS Editor
Copyright (C) Microsoft Corporation, 1987-1991.
All rights reserved.
Press Enter to see the Survival Guide.
c:\Press ESC to close this starting menu.

Oracle VM VirtualBox Manager

File Machine Help

New Settings Discard Show

MikroTik_RouterOS_CHR
2.6 Powered Off

DOS Restoring

Linux 2.6 Powered Off

kali-linux-2022.4-virtualb...
2.6 Powered Off

General
Name: DOS
Operating System: DOS

System

Base Memory: 32 MB

Boot Order: Floppy, Optical, Hard Disk

Acceleration: VT-x/AMD-V, Nested Paging

Display

Video Memory: 9 MB

Scale-factor: 1.50

Graphics Controller: VBoxVGA

Remote Desktop Server: Disabled

Recording: Disabled

Storage

Audio

Host Driver: Windows DirectSound

Controller: SoundBlaster 16

Network

Adapter 1: PCnet-FAST III (NAT)

USB

Disabled

Shared

You have the Auto capture keyboard option turned on. This will cause the Virtual Machine to automatically capture the keyboard every time the VM window is activated and make it unavailable to other windows.

Untitled

Help

File Machine View Input Devices Help

Welcome to the MS-DOS Editor

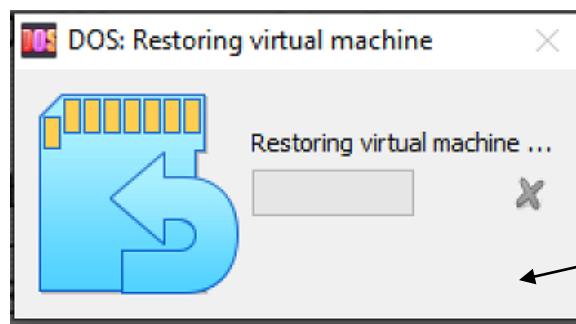
Copyright (C) Microsoft Corporation, 1987-1991.

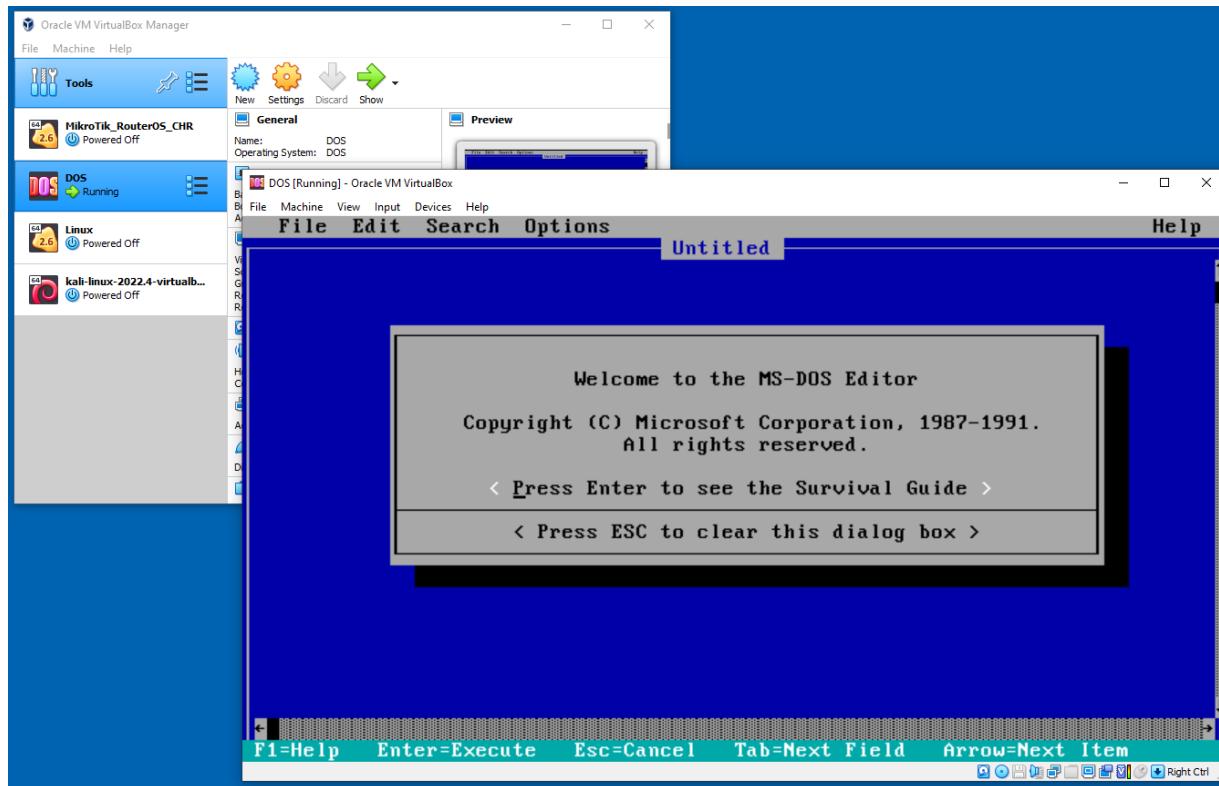
All rights reserved.

Press Enter to see the Survival Guide.

c:\Press ESC to close this starting menu.

Restoring virtual machine ...





Implementation of Processes

Process management	Memory management	File management
Registers Program counter Program status word Stack pointer Process state Priority Scheduling parameters Process ID Parent process Process group Signals Time when process started CPU time used Children's CPU time Time of next alarm	Pointer to text segment info Pointer to data segment info Pointer to stack segment info	Root directory Working directory File descriptors User ID Group ID

Some of the fields of a typical process table entry.

Implementation of Processes

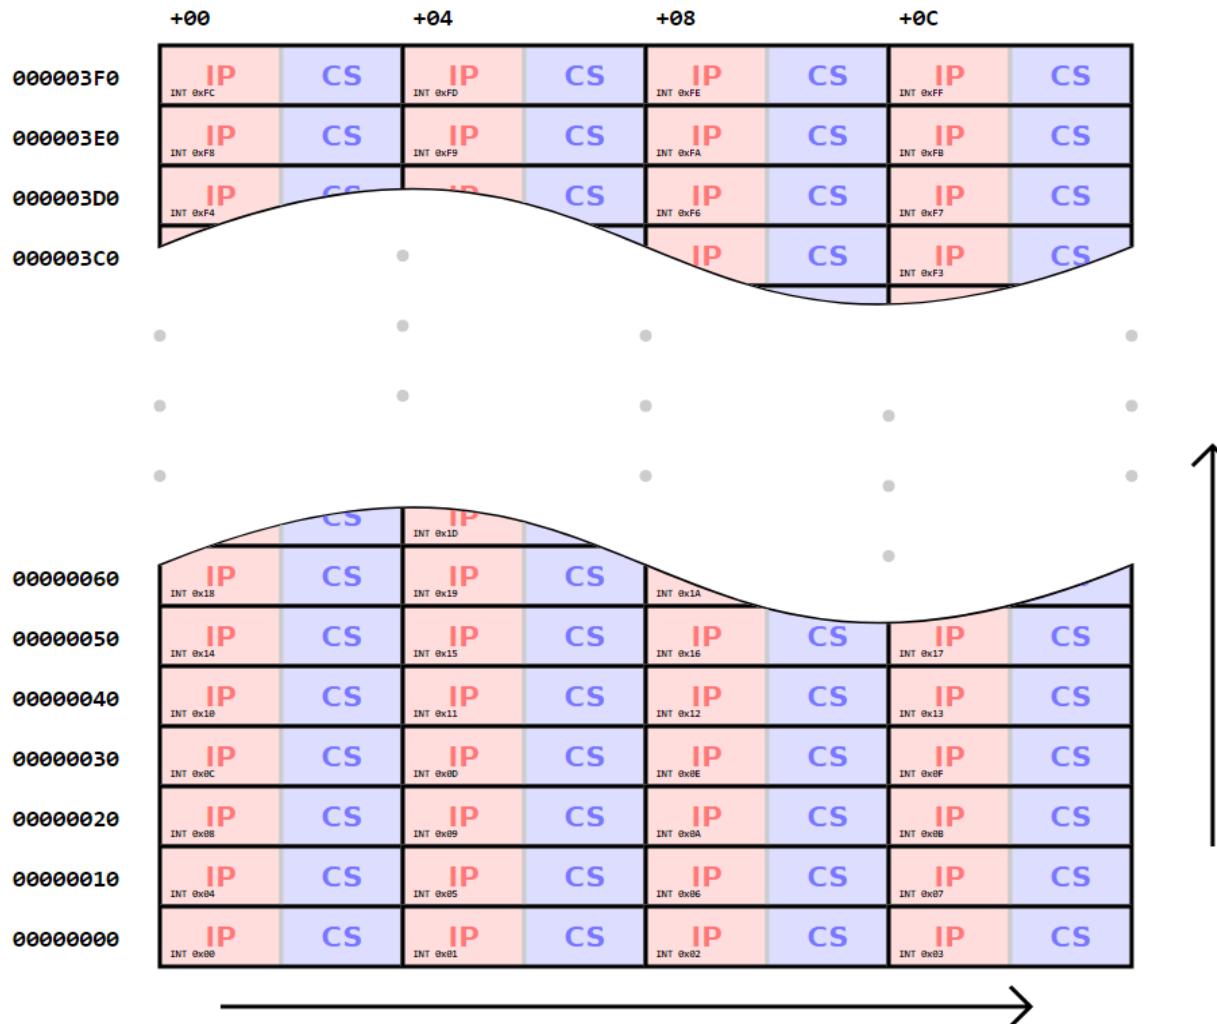
The illusion of multiple sequential processes maintained on one (or each) CPU. Suppose that user process is running when a disk interrupt happens.

Skeleton of what the lowest level of the operating system does when an interrupt occurs.

1. Hardware stacks program counter, etc.
2. Hardware loads new program counter from interrupt vector.
3. Assembly language procedure saves registers.
4. Assembly language procedure sets up new stack.
5. C interrupt service runs (typically reads and buffers input).
6. Scheduler decides which process is to run next.
7. C procedure returns to the assembly code.
8. Assembly language procedure starts up new current process.

Associated with each I/O class is a location (typically at a fixed location near the bottom of memory) called the **interrupt vector**.

Interrupt Vector Table



Interrupt Vector Table of x86 processors running in real mode.

Source: WIKIPEDIA

Modeling Multiprogramming

CPU utilization a probabilistic viewpoint

Suppose that a process spends a fraction p of its time waiting for I/O to complete. With n processes in memory at once, the probability that all n processes are waiting for I/O (in which case the CPU will be idle) is p^n . The CPU utilization is then given by the formula:

$$\text{CPU utilization} = 1 - p^n$$

Modeling Multiprogramming

	P1	P2	P3
0 idle	1-p	1-p	1-p
1 idle	p 1-p 1-p	1-p p 1-p	1-p 1-p p
2 idle	1-p p p	p 1-p p	p p 1-p
3 idle	p	p	p

$$(1 - p)^3$$

$$p(1 - p)^2$$

$$p(1 - p)^2 \quad 3p(1 - p)^2$$

$$p(1 - p)^2$$

$$p^2(1 - p)$$

$$p^2(1 - p) \quad 3p^2(1 - p)$$

$$p^2(1 - p)$$

$$p^3$$

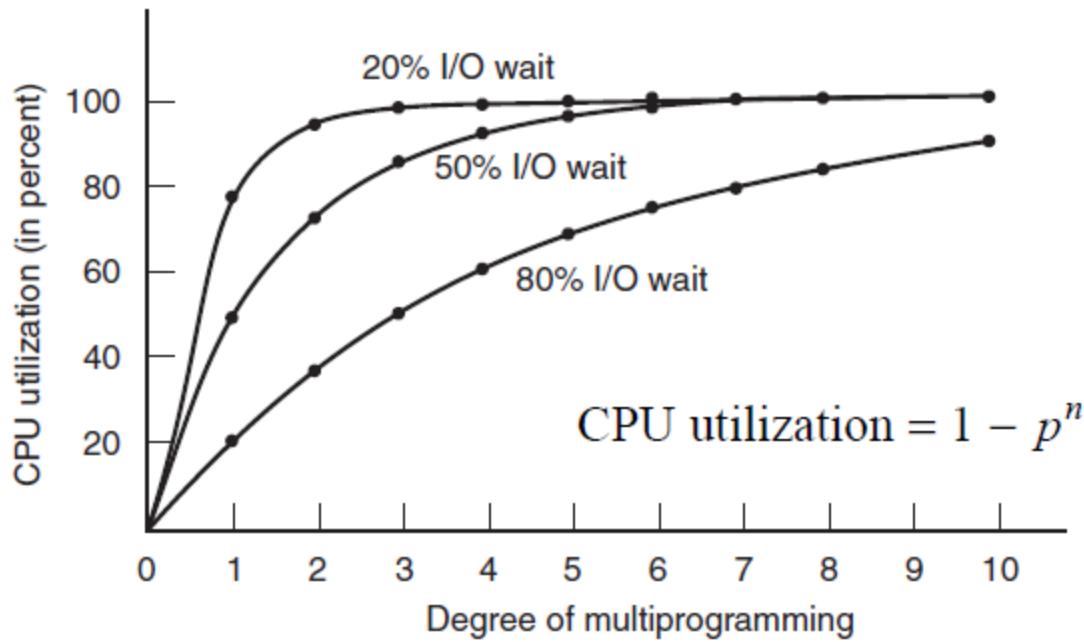
CPU - utilization

$$U = (1 - p)^3 + 3p(1 - p)^2 + 3p^2(1 - p)$$

$$U = 1 - p^3$$

Modeling Multiprogramming

CPU utilization
a probabilistic viewpoint



Modeling Multiprogramming

CPU utilization a probabilistic viewpoint

From the figure it is clear that if processes spend 80% of their time waiting for I/O, at least 10 processes must be in memory at once to get the CPU waste below 10%.

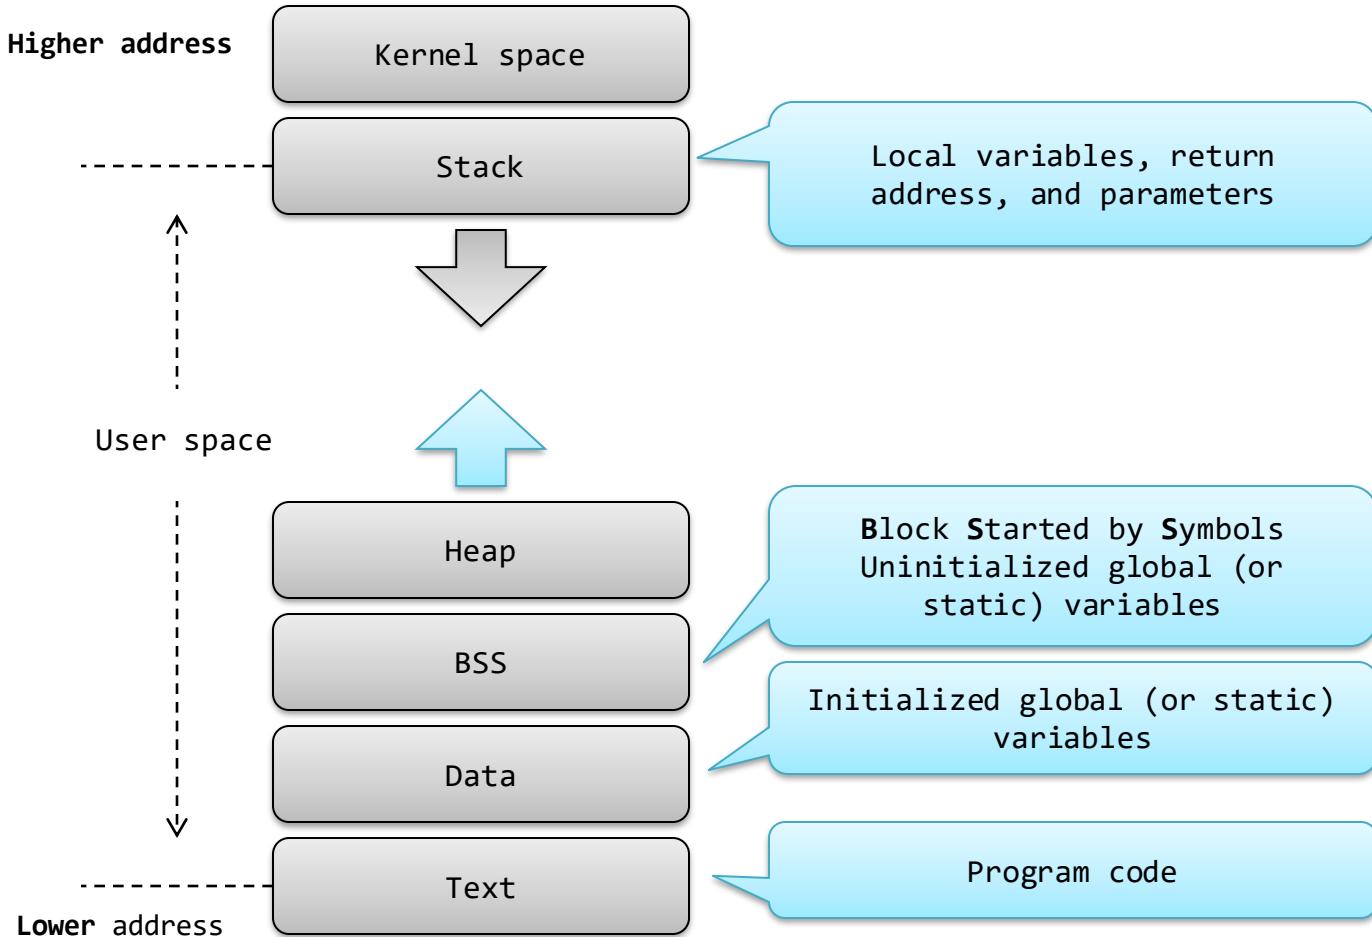
$$\text{CPU waste} = (0.8)^{10} \approx 0.1074$$

When you realize that an interactive process waiting for a user to type something at a terminal (or click on an icon) is in I/O wait state, it should be clear that I/O wait times of 80% and more are not unusual.

It should be pointed out that the probabilistic model assumes that all n processes are independent.

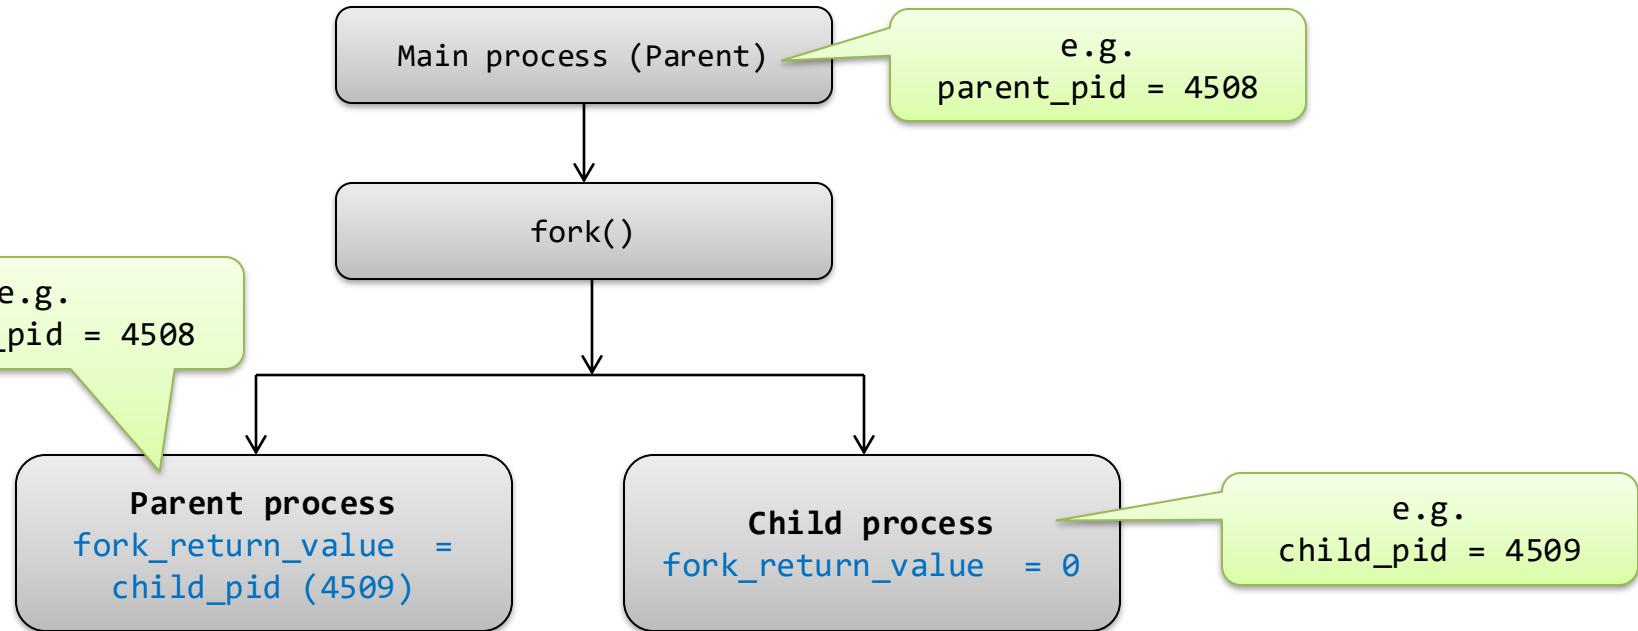
Linux

Processes in Unix



Linux

Processes creation



Linux

Processes creation

```
#define TRUE 1

while (TRUE) {
    type_prompt( );
    read_command(command, parameters);
    /* repeat forever */
    /* display prompt on the screen */
    /* read input from terminal */

    if (fork( ) != 0) {
        /* Parent code. */
        waitpid(-1, &status, 0);
        /* fork off child process */
        /* wait for child to exit */
    } else {
        /* Child code. */
        execve(command, parameters, 0);
        /* execute command */
    }
}
```

A stripped-down shell. Throughout this book, *TRUE* is assumed to be defined as 1.

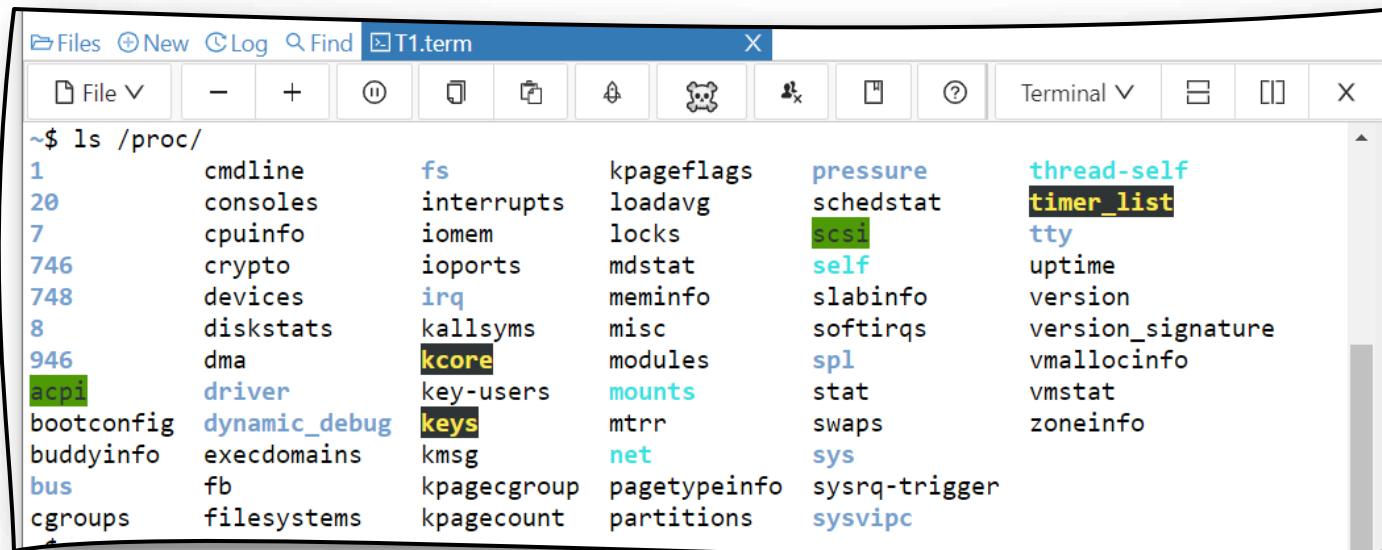
Linux

process information pseudo-filesystem (proc)

man proc

The proc filesystem is a pseudo-filesystem which provides an interface to kernel data structures. It is commonly mounted at /proc.

The /proc/ directory contains a wealth of information detailing system hardware and any running processes. In addition, some of the files within /proc/ can be manipulated by users and applications to communicate configuration changes to the kernel([fedora documentation](#)).



```
File New Log Find T1.term X
File - + ⌂ ⌄ ⌅ ⌆ ⌇ ⌈ ⌉ ⌊ ⌋ ⌃ Terminal ⌄ ⌁ ⌂ ⌃ X
~$ ls /proc/
1      cmdline      fs          kpageflags   pressure    thread-self
20     consoles     interrupts  loadavg      schedstat  timer_list
7      cpuinfo      iomem       locks        scsi        tty
746    crypto       ioports     mdstat      self        uptime
748    devices      irq        meminfo     slabinfo   version
8      diskstats   kallsyms   misc        softirqs   version_signature
946    dma         kcore      modules     spl        vmallocinfo
acpi   driver      key-users  mounts     stat        vmstat
bootconfig dynamic_debug keys      mtrr       swaps     zoneinfo
buddyinfo  execdomains kmsg      net        sys
bus      fb          kpagecgrou pagetypeinfo sysrq-trigger
cgroups   filesystems kpgecount partitions sysvipc
```

Linux

proc - process information pseudo-filesystem

Every process in Linux is represented by a directory in /proc e.g.

```
New process  
-> ~$ sleep 10000 &  
[1] 1680  
~$ ls /proc/  
1 diskstats kpagecount softirqs  
1680 dma kpageflags spl  
169 driver loadavg stat  
2 dynamic_debug locks swaps  
3 ins mdstat sys  
4 filesystems meminfo sysrq-trigger  
5 misc sysvipc  
748
```

Process ID

New folder is created

The screenshot shows a terminal window with the following content:

```
~$ sleep 10000 &
[1] 1680
~$ ls /proc/
1 diskstats kpagecount softirqs
1680 dma kpageflags spl
169 driver loadavg stat
2 dynamic_debug locks swaps
3 ins mdstat sys
4 filesystems meminfo sysrq-trigger
5 misc sysvipc
748
```

Annotations with callout boxes:

- A box labeled "New process" points to the output of the sleep command, specifically the PID [1] 1680.
- A box labeled "Process ID" points to the PID 1680 in the /proc listing.
- A box labeled "New folder is created" points to the newly created directory 1680 in the /proc listing.

Linux

proc - process information pseudo-filesystem

```
~$ sleep 10000 &
[1] 1680
~$ ls /proc/1680/
arch_status          fd          numa_maps      smaps_rollup
attr                fdinfo      oom_adj        stack
autogroup           gid_map    oom_score      stat
auxv               io          oom_score_adj  statm
cgroup              limits     pagemap       status
clear_refs          loginuid   patch_state   syscall
cmdline             map_files  personality   task
comm                maps       projid_map   timens_offsets
coredump_filter     mem        root          timers
cpu_resctrl_groups mountinfo   sched         timerslack_ns
cpuset              mounts     schedstat    uid_map
cwd                 mountstats sessionid    wchan
environ             net        setgroups   smaps
exe
~$
```

Linux

proc - process information pseudo-filesystem

e.g. process address space

```
~$ sleep 10000 &
[1] 1680
~$ cat /proc/1680/maps
5558f1db2000-5558f1db4000 r--p 00000000 00:cf 26481965 /usr/bin/sleep
5558f1db4000-5558f1db8000 r-xp 00002000 00:cf 26481965 /usr/bin/sleep
5558f1db8000-5558f1dba000 r--p 00006000 00:cf 26481965 /usr/bin/sleep
5558f1dbb000-5558f1dbc000 r--p 00008000 00:cf 26481965 /usr/bin/sleep
5558f1dbc000-5558f1dbd000 rw-p 00009000 00:cf 26481965 /usr/bin/sleep
5558f326d000-5558f328e000 rw-p 00000000 00:00 0 [heap]
7f31af675000-7f31af6a7000 r--p 00000000 00:cf 41908341 /usr/lib/locale/C.UTF-8/LC_CTYPE
7f31af6a7000-7f31af81a000 r--p 00000000 00:cf 41908340 /usr/lib/locale/C.UTF-8/LC_COLLATE
7f31af81a000-7f31af83c000 r--p 00000000 00:cf 41899982 /lib/x86_64-linux-gnu/libc-2.31.so
7f31af83c000-7f31af9b4000 r-xp 00022000 00:cf 41899982 /lib/x86_64-linux-gnu/libc-2.31.so
7f31af9b4000-7f31afa02000 r--p 0019a000 00:cf 41899982 /lib/x86_64-linux-gnu/libc-2.31.so
7f31afa02000-7f31afa06000 r--p 001e7000 00:cf 41899982 /lib/x86_64-linux-gnu/libc-2.31.so
7f31afa06000-7f31afa08000 rw-p 001eb000 00:cf 41899982 /lib/x86_64-linux-gnu/libc-2.31.so
7f31afa08000-7f31afa0e000 rw-p 00000000 00:00 0
7f31afa3a000-7f31afa3b000 r--p 00000000 00:cf 41908346 /usr/lib/locale/C.UTF-8/LC_NUMERIC
7f31afa3b000-7f31afa3c000 r--p 00000000 00:cf 41908349 /usr/lib/locale/C.UTF-8/LC_TIME
```

Linux

process monitoring tools

Pstree, ps, and top(htop)

`pstree` - display a tree of processes

e.g.

pstree -p -T

`ps` - report a snapshot of the current processes

e.g.

ps aux

`top` - display Linux processes

e.g.

top

Linux

process monitoring tools

Pstree, ps, and top(htop)

```
~$  
~$ pstree -p -T  
tini(1)—sh(7)—node(8)—bash(746)—pstree(859)  
          |           |  
          |           bash(748)  
          |           |  
          |           sshd(20)  
~$
```

```
~$  
~$ ps ax  
 PID TTY      STAT   TIME COMMAND  
   1 ?        Ss      0:00 /cocalc/bin/tini -v -g -- sh -c env -:  
   7 ?        SN      0:00 sh -c env -i /cocalc/init/init.sh $COO  
   8 ?        RN1     0:10 node --optimize-for-size --gc-interval  
  20 ?        SN      0:00 sshd: /usr/sbin/sshd -D -p 2222 -h /tr  
 314 pts/0    SNs     0:00 /bin/bash  
 401 pts/0    SN      0:00 sleep 10000  
 403 pts/0    RN+    0:00 ps ax  
~$  
~$
```

Linux

process monitoring tools

Pstree, ps, and top(htop)

Tasks: 160 total, 1 running, 159 sleeping, 0 stopped, 0 zombie											
%Cpu(s): 0.7 us, 0.0 sy, 0.0 ni, 99.3 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st											
MiB Mem : 971.5 total, 117.5 free, 402.5 used, 451.6 buff/cache											
MiB Swap: 448.5 total, 307.1 free, 141.4 used. 417.1 avail Mem											
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1233	root	20	0	306476	69736	13224	S	0.7	7.0	0:06.99	Xorg
1935	asp	20	0	313944	27712	19276	S	0.7	2.8	0:02.74	xfce4-terminal
8140	asp	20	0	11856	3604	3096	R	0.3	0.4	0:00.06	top
1	root	20	0	102960	10488	6744	S	0.0	1.1	0:02.40	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd
3	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	rcu_gp
4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	rcu_par_gp
6	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/0:0H-events_highpri
7	root	20	0	0	0	0	I	0.0	0.0	0:01.74	kworker/0:1-events
9	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	mm_percpu_wq
10	root	20	0	0	0	0	S	0.0	0.0	0:00.00	rcu_tasks_rude
11	root	20	0	0	0	0	S	0.0	0.0	0:00.00	rcu_tasks_trace
12	root	20	0	0	0	0	S	0.0	0.0	0:00.31	ksoftirqd/0
13	root	20	0	0	0	0	I	0.0	0.0	0:00.64	rcu_sched
14	root	rt	0	0	0	0	S	0.0	0.0	0:00.03	migration/0

Linux

process monitoring tools

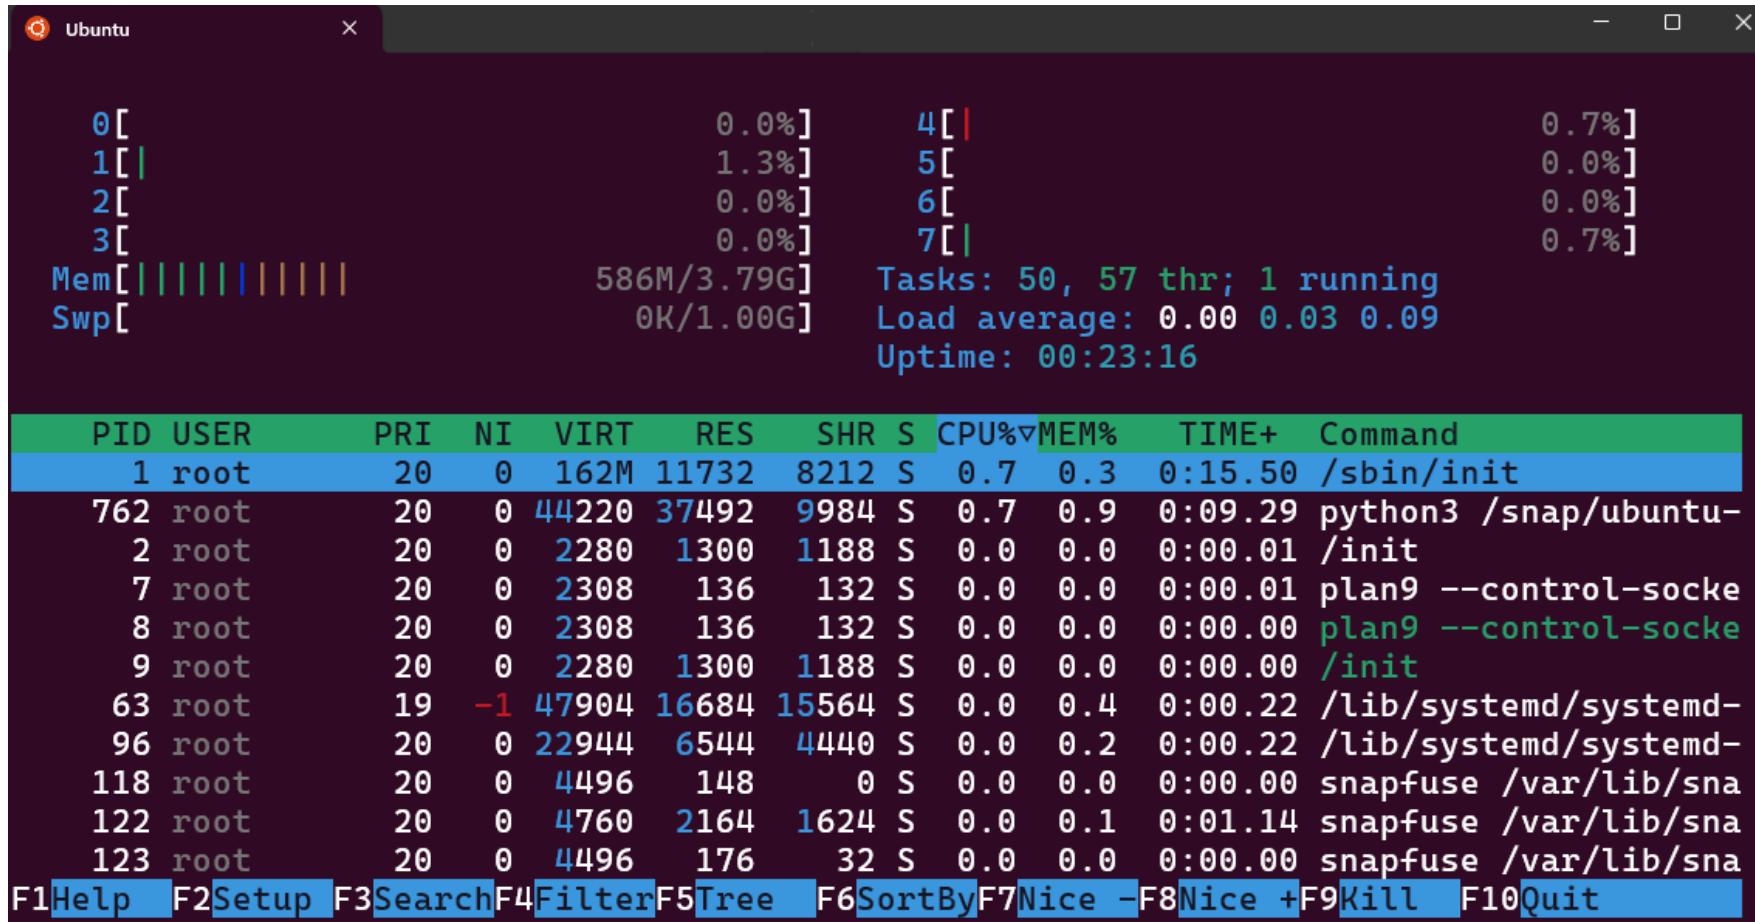
Pstree, ps, and top(htop)

Tasks: 160 total, 1 running, 159 sleeping, 0 stopped, 0 zombie										
%Cpu(s): 0.7 us, 0.0 sy, 0.0 ni, 99.3 id, 0.0 wa, 0.0 st										
Mem: 971.5 total, 117.5 free, 402.5 used, 741.6 buff/cache										
Process ID	User	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+ COMMAND
1233	root	20	0	306476	69736	13224	S	0.7	7.0	0:06.99 Xorg
1935	as	20	0	313944	7112	19276	S	0.7	2.8	0:02.74 xfce4-terminal
8	User name	Nice value		1856	Resident size		R	0.3	0.4	0:00.06 top
1	root	20	0	102960	10488	674	S	0.0	0.1	0:02.40 systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00 kthreadd
3	root	0	-20	0	0	0	I	0.0	0.0	rcu_gp
4	root	0	-20	0	0	0	I	0.0	0.0	rcu_par_gp
6	root	0	-20	0	0	0	I	0.0	0.0	kworker/0:0H-events_highpri
7	root	20	0	0	0	0	I	0.0	0.0	kworker/0:1-events
9	root	0	-20	0	0	0	I	0.0	0.0	mm_percpu_wq
10	root	20	0	0	0	0	S	0.0	0.0	0:00.00 rcu_tasks_rude
11	root	20	0	0	0	0	S	0.0	0.0	0:00.00 rcu_tasks_trace
12	root	20	0	0	0	0	S	0.0	0.0	0:00.31 ksoftirqd/0
13	root	20	0	0	0	0	I	0.0	0.0	0:00.64 rcu_sched
14	root	rt	0	0	0	0	S	0.0	0.0	0:00.03 migration/0

Linux

process monitoring tools

Pstree, ps, and top(htop)



Linux

Process state

man ps

- D uninterruptible sleep (usually IO and cannot be killed)
- I Idle kernel thread
- R running or runnable (on run queue)
- S interruptible sleep (waiting for an event to complete and can be killed)
- T stopped by job control signal
- X dead (should never be seen)
- Z defunct ("zombie") process, terminated but not reaped by its parent

- < high-priority (not nice to other users)
- N low-priority (nice to other users)
- + is in the foreground process group

Linux

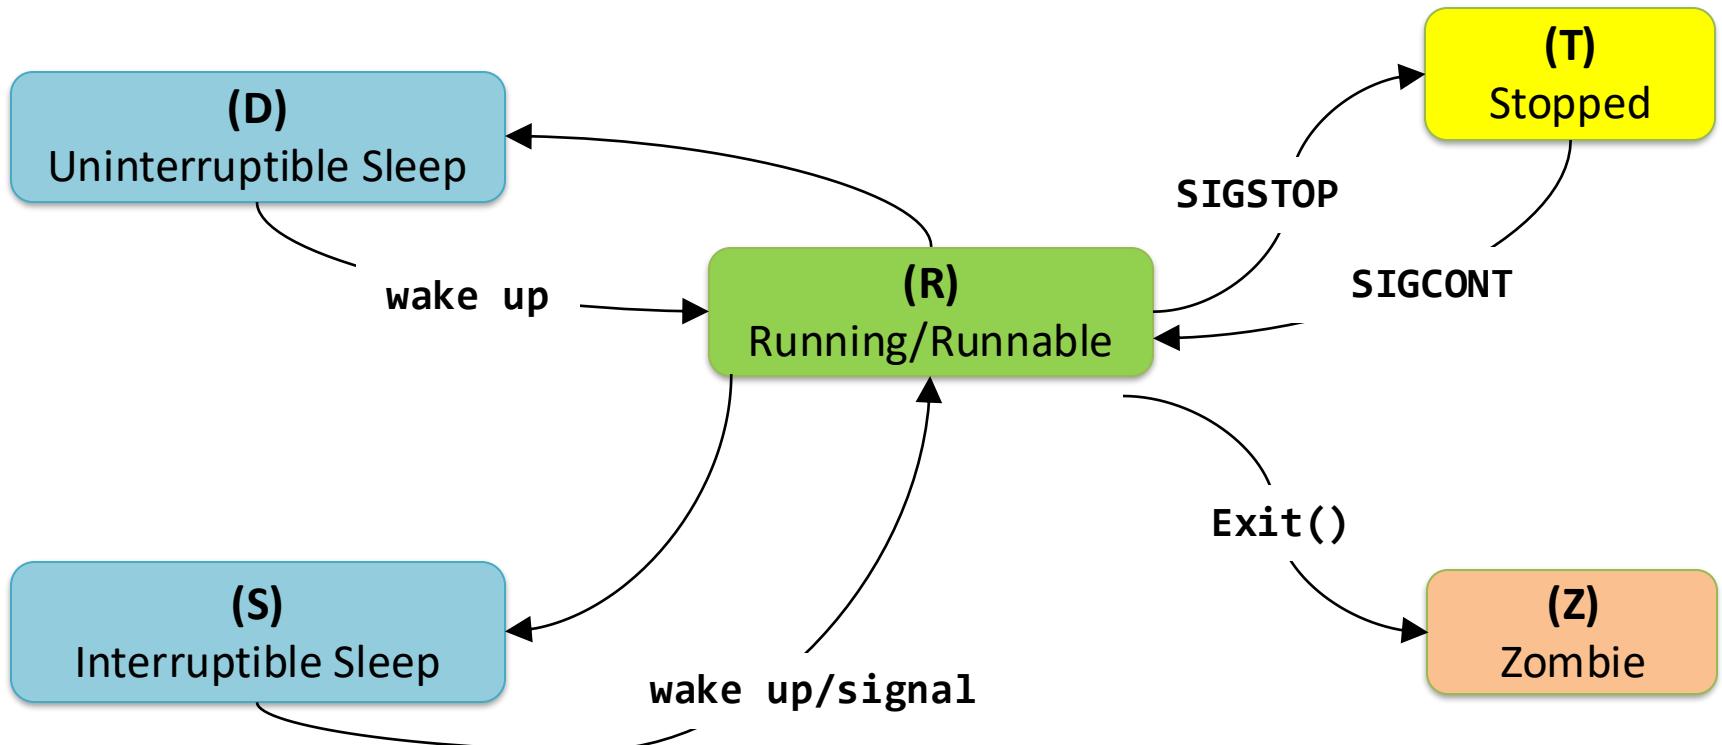
Process state

S interruptible sleep (waiting for an event to complete)

```
~$  
~$ ps ax  
 PID TTY      STAT   TIME COMMAND  
   1 ?        Ss     0:00 /cocalc/bin/tini -v -g -- sh -c env -i  
   7 ?        SN     0:00 sh -c env -i /cocalc/init/init.sh $CO  
   8 ?        RN1    0:10 node --optimize-for-size --gc-interval  
  20 ?        SN     0:00 sshd: /usr/sbin/sshd -D -p 2222 -h /tm  
 314 pts/0    SNs    0:00 /bin/bash  
 401 pts/0    SN     0:00 sleep 10000  
 403 pts/0    RN+    0:00 ps ax  
~$
```

Linux

Process state



Linux

Process state

T stopped by job control signal

```
asp@asp-VirtualBox:~/Desktop/OS$ ./a.out
[1]+  Stopped                  ./a.out
asp@asp-VirtualBox:~/Desktop/OS$
```

```
Tasks: 161 total,  1 running, 159 sleeping,  1 stopped,  0 zombie
%Cpu(s): 0.3 us, 0.0 sy, 0.0 ni, 99.7 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 971.5 total,   87.0 free, 446.5 used, 438.1 buff/cache
MiB Swap: 448.5 total, 307.4 free, 141.1 used. 367.3 avail Mem
```

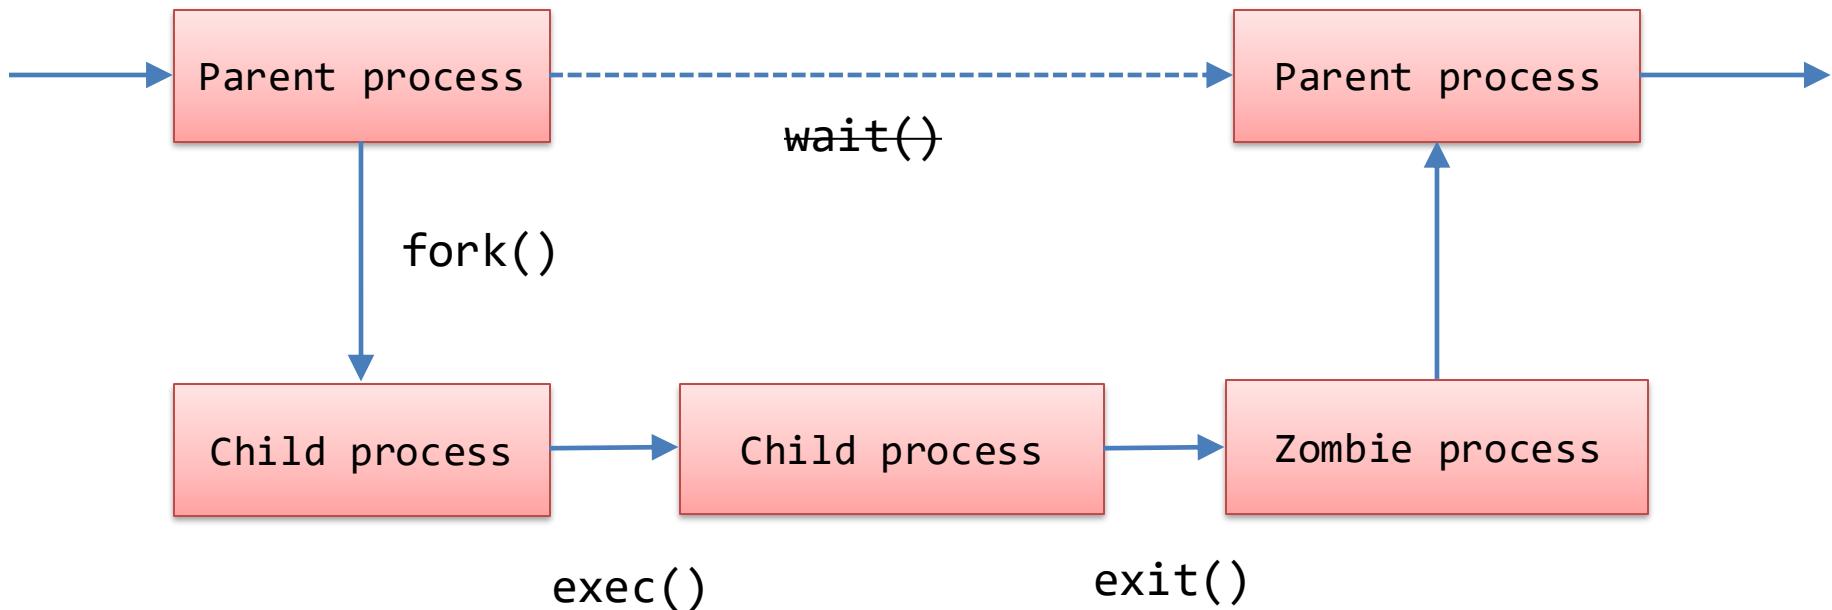
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
8266	asp	20	0	2364	512	448	T	0.0	0.1	0:00.00	a.out

Linux

Zombie process

The parent process is responsible for waiting for child process to terminate and then cleaning up the child process entry from the process table.

The zombie process can be easily created when the parent process does not wait for the child process termination. A zombie process doesn't consume many system resources, but it does continue to live in the system's task list.



Linux

Zombie process

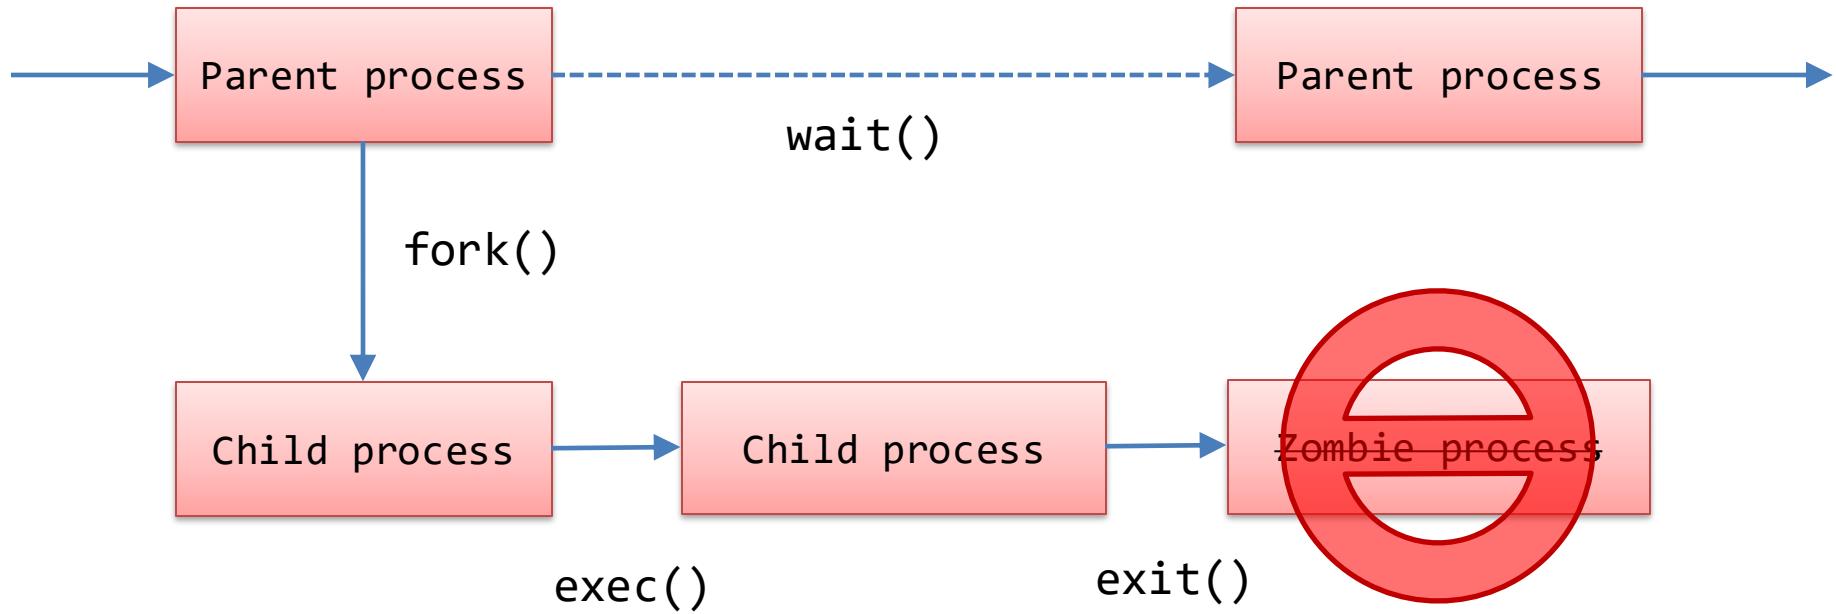
```
~$ ./a.out
Child: pid is 1550
Parent: pid is 1549
~$
```

```
top - 23:58:38 up 4:15, 0 users, load average: 1.84, 1.99, 2.58
Tasks: 9 total, 1 running, 7 sleeping, 0 stopped, 1 zombie
%Cpu(s): 13.2 us, 6.0 sy, 1.8 ni, 75.4 id, 2.7 wa, 0.0 hi, 0.8 si, 0.0 st
MiB Mem : 32104.4 total, 10043.8 free, 3500.7 used, 18559.8 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used. 26939.4 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
8	user	38	18	963500	111260	33800	S	2.0	0.3	1:27.32	node
1	user	20	0	2500	584	520	S	0.0	0.0	0:00.06	tini
7	user	38	18	2616	600	528	S	0.0	0.0	0:00.00	sh
20	user	38	18	12180	6788	5956	S	0.0	0.0	0:00.01	sshd
363	user	38	18	8116	6056	3308	S	0.0	0.0	0:00.13	bash
910	user	38	18	8116	6232	3484	S	0.0	0.0	0:00.08	bash
1549	user	38	18	2364	512	444	S	0.0	0.0	0:00.00	a.out
1550	user	38	18	0	0	0	Z	0.0	0.0	0:00.00	a.out
1551	user	38	18	7904	3612	3096	R	0.0	0.0	0:00.00	top

Linux

Zombie process



Linux

Zombie process

```
~$ ./a.out
Child: pid is 1850
Parent: pid is 1849
~$
```

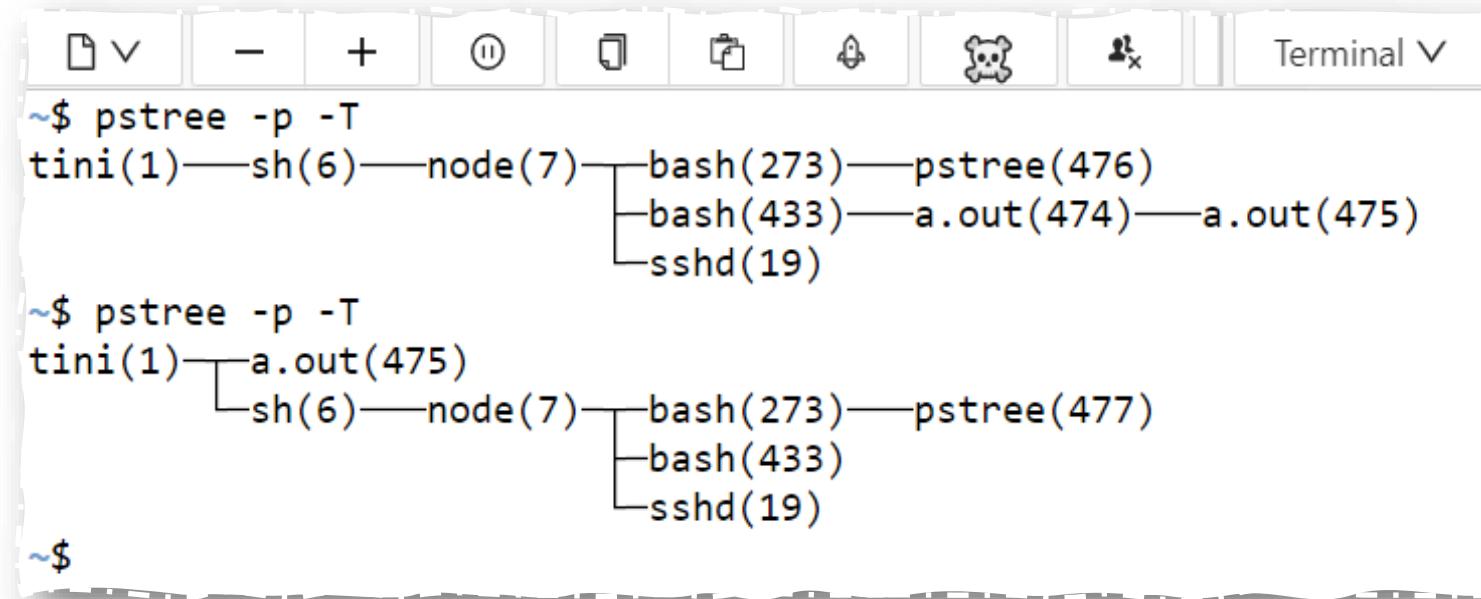
```
top - 00:14:44 up 4:31, 0 users, load average: 1.55, 1.86, 2.08
Tasks: 8 total, 1 running, 7 sleeping, 0 stopped, 0 zombie
%Cpu(s): 11.3 us, 4.4 sy, 1.8 ni, 81.3 id, 0.7 wa, 0.0 hi, 0.6 si, 0.0 st
MiB Mem : 32104.4 total, 10316.8 free, 3168.3 used, 18619.2 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used. 27271.3 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
8	user	38	18	964264	113720	33800	S	10.6	0.3	2:11.88	node
1	user	20	0	2500	584	520	S	0.0	0.0	0:00.08	tini
7	user	38	18	2616	600	528	S	0.0	0.0	0:00.00	sh
20	user	38	18	12180	6788	5956	S	0.0	0.0	0:00.01	sshd
363	user	38	18	8116	6056	3308	S	0.0	0.0	0:00.14	bash
910	user	38	18	8116	6232	3484	S	0.0	0.0	0:00.08	bash
1848	user	38	18	7904	3540	3024	R	0.0	0.0	0:00.00	top
1849	user	38	18	2364	572	508	S	0.0	0.0	0:00.00	a.out

Linux

Orphan process

The parent dies (or terminates) before its child. It means that the child process is still running even though the parent process has terminated. The child process will be automatically adopted by the **init** process (the parent of all the processes with the pid of 1).



A screenshot of a Linux desktop environment showing a terminal window. The terminal window title is "Terminal". The terminal displays two outputs of the `pstree` command.

```
~$ pstree -p -T
tini(1)—sh(6)—node(7)—bash(273)—pstree(476)
                                |           |
                                |           bash(433)—a.out(474)—a.out(475)
                                |           |
                                |           sshd(19)

~$ pstree -p -T
tini(1)—a.out(475)
                  |           |
                  sh(6)—node(7)—bash(273)—pstree(477)
                                |           |
                                |           bash(433)
                                |           |
                                |           sshd(19)

~$
```

The first `pstree` output shows a process tree starting from `tini(1)`. It branches through `sh(6)`, `node(7)`, `bash(273)`, and `pstree(476)`. From `bash(273)`, it further branches to `bash(433)`, `a.out(474)`, and `a.out(475)`. Finally, `a.out(475)` has a child process `sshd(19)`.

The second `pstree` output shows a similar process tree. It starts from `tini(1)` and branches to `a.out(475)`. `a.out(475)` then has a child process `sh(6)`. `sh(6)` further branches to `node(7)`, `bash(273)`, and `pstree(477)`. From `bash(273)`, it branches to `bash(433)` and `sshd(19)`.

Linux

Orphan process

